Chapter 5. Boosting Methods

Predictive Modeling, Lecture 4 (AdaBoost, gradient boosting, XGBoost, LightGBM, CatBoost)

Contents

1 From Averaging to Sequential Correction **2**

2 AdaBoost **2**

2.1 Algorithm 2

2.2 AdaBoost is forward stagewise fitting of the exponential loss 2

3 Gradient Boosting Machines **3**

3.1 Algorithm 3

3.2 Loss functions and their pseudo-residuals 4

3.3 Shrinkage, subsampling and early stopping 4

3.4 Limitations of the standard GBM 4

4 XGBoost: Regularised Second-Order Boosting **4**

5 LightGBM **5**

6 CatBoost **6**

6.1 The curse of dimensionality of categorical variables 6

6.2 Naive target statistics leak the target 7

7 Summary and Practical Rules **7**

1 From Averaging to Sequential Correction

Random forest (Ch. 4)	Gradient boosting (this chapter)
$\hat{f}_{\mathrm{RF}}(x) = \frac{1}{B}\sum_{b=1}^{B} T_b(x)$	$F_M(x) = F_0(x) + \sum_{m=1}^{M} \eta\, h_m(x)$
trees learned **independently** on different bootstrap samples	h_m: a shallow tree learning the **error (residual) direction** left by the previous model
averaging reduces the instability of individual trees	η: the **learning rate**, shrinking each correction to prevent overfitting
every tree looks at the same target Y	small corrections ηh_m accumulate **sequentially**, not independently

Random forests mainly reduce variance; boosting mainly reduces bias by adding many weak, shallow learners. This chapter covers *by which principle* boosting learns on tabular data and *under which conditions* to choose it.

2 AdaBoost

2.1 Algorithm

Algorithm: AdaBoost (labels $y_i \in \{-1, +1\}$)

Initialise $D_1(i) = 1/N$. For $b = 1, \ldots, B$:

1. **Train**: fit a classifier $h_b(x) \in \{-1, +1\}$ using weights D_b.
2. **Evaluate**: weighted error $\epsilon_b = \sum_{i=1}^{N} D_b(i)\, \mathbb{I}\{h_b(x_i) \neq y_i\}$.
3. **Weight**: $\alpha_b = \frac{1}{2} \log \frac{1-\epsilon_b}{\epsilon_b}$.
4. **Update**: $D_{b+1}(i) = \frac{D_b(i) \exp\{-\alpha_b y_i h_b(x_i)\}}{Z_b}$, Z_b the normaliser.

Output: $H(x) = \mathrm{sign}\Big(\sum_{b=1}^{B} \alpha_b h_b(x)\Big)$.

Core mechanism: the more an observation was misclassified by previous models, the larger its weight D_{b+1} in the next round ($\exp(+\alpha_b)$ for errors, $\exp(-\alpha_b)$ for correct cases); the better a model (smaller ϵ_b), the larger its say α_b in the final vote. A learner with $\epsilon_b = \frac{1}{2}$ gets $\alpha_b = 0$; one with $\epsilon_b > \frac{1}{2}$ would get a negative weight (it is flipped).

2.2 AdaBoost is forward stagewise fitting of the exponential loss

Theorem: derivation of α_b and the weight update

Let $F_{b-1} = \sum_{s<b} \alpha_s h_s$ and consider the exponential loss $L(y, F) = e^{-yF}$. Choosing (α, h) to minimise $\sum_i \exp\{-y_i(F_{b-1}(x_i) + \alpha h(x_i))\}$ gives: h_b minimises the weighted error with weights $D_b(i) \propto e^{-y_i F_{b-1}(x_i)}$, and $\alpha_b = \frac{1}{2} \log \frac{1-\epsilon_b}{\epsilon_b}$. The weights evolve exactly by the AdaBoost update.

Proof. With $w_i = e^{-y_i F_{b-1}(x_i)}$, normalised to $D_b(i)$, and $y_i h(x_i) = \pm 1$,

$$\sum_i D_b(i) e^{-\alpha y_i h(x_i)} = e^{-\alpha}(1-\epsilon) + e^{\alpha}\epsilon = e^{-\alpha} + (e^{\alpha} - e^{-\alpha})\epsilon.$$

For any $\alpha > 0$ this is increasing in the weighted error ϵ, so h_b is the weighted-error minimiser. Differentiating in α: $-e^{-\alpha}(1-\epsilon) + e^{\alpha}\epsilon = 0 \iff e^{2\alpha} = (1-\epsilon)/\epsilon$. Finally $w_i^{\text{new}} = w_i e^{-\alpha_b y_i h_b(x_i)}$, which after normalisation is the update rule. ■

Theorem: training-error bound

$Z_b = 2\sqrt{\epsilon_b(1-\epsilon_b)}$ and

$$\frac{1}{N}\sum_{i=1}^{N} \mathbb{I}\{H(x_i) \neq y_i\} \leq \prod_{b=1}^{B} Z_b \leq \exp\Big(-2\sum_{b=1}^{B}\gamma_b^2\Big), \qquad \gamma_b = \tfrac{1}{2} - \epsilon_b.$$

Proof. Plugging α_b into $Z_b = e^{-\alpha_b}(1-\epsilon_b) + e^{\alpha_b}\epsilon_b$ gives $2\sqrt{\epsilon_b(1-\epsilon_b)}$. Unrolling the update, $D_{B+1}(i) = \frac{e^{-y_i F_B(x_i)}}{N\prod_b Z_b}$, and these sum to 1, so $\frac{1}{N}\sum_i e^{-y_i F_B(x_i)} = \prod_b Z_b$. Since $\mathbb{I}\{H(x_i) \neq y_i\} \leq e^{-y_i F_B(x_i)}$, the first inequality follows. Finally $2\sqrt{\epsilon(1-\epsilon)} = \sqrt{1-4\gamma^2} \leq e^{-2\gamma^2}$. ■

So if every weak learner is slightly better than chance ($\gamma_b \geq \gamma > 0$) the training error falls exponentially fast. The population minimiser of the exponential loss is $f^\star(x) = \frac{1}{2}\log\frac{P(Y=1|x)}{P(Y=-1|x)}$ (minimise $pe^{-f} + (1-p)e^{f}$), so AdaBoost estimates half the log-odds.

Example: AdaBoost on a toy case (classic illustration)

Ten points in the plane with decision stumps. Round 1: a vertical stump misclassifies three + points, $\epsilon_1 = 0.30$, $\alpha_1 = \frac{1}{2}\log(0.7/0.3) = 0.42$; the three misclassified points are enlarged in D_2. Round 2: another vertical stump, $\epsilon_2 = 0.21$, $\alpha_2 \approx 0.65$; three − points it misclassifies grow in D_3. Round 3: a horizontal stump, $\epsilon_3 = 0.14$, $\alpha_3 \approx 0.92$. The final classifier

$$H_{\text{final}} = \text{sign}(0.42\, h_1 + 0.65\, h_2 + 0.92\, h_3)$$

is a staircase boundary that no single stump can represent. (With the formula, 0.21 and 0.14 give 0.66 and 0.91; the figures round the weighted errors.)

3 Gradient Boosting Machines

3.1 Algorithm

Algorithm: GBM

Initialise $F_0(x) = \arg\min_\gamma \sum_{i=1}^{N} L(y_i, \gamma)$. For $b = 1, \ldots, B$:

1. **Gradient**: pseudo-residuals $r_{ib} = -\Big[\dfrac{\partial L(y_i, F(x_i))}{\partial F(x_i)}\Big]_{F=F_{b-1}}$.
2. **Fit**: fit a regression tree $h_b(x)$ (leaves R_{jb}) to the targets r_{ib}; optionally set each leaf value by a line search $\gamma_{jb} = \arg\min_\gamma \sum_{x_i \in R_{jb}} L(y_i, F_{b-1}(x_i) + \gamma)$.
3. **Update**: $F_b(x) = F_{b-1}(x) + \eta\, h_b(x)$.

Output: $F_B(x) = F_0(x) + \sum_{b=1}^{B} \eta\, h_b(x)$.

Core mechanism: if the loss L is squared error, the gradient r_{ib} coincides exactly with the ordinary residual $y_i - F_{b-1}(x_i)$. Unlike AdaBoost, **any differentiable loss** can be used. It is the principle of **gradient descent in function space**: the vector $(F(x_1), \ldots, F(x_N))$ is moved in the negative gradient direction, and the tree h_b is the best approximation of that direction that can also be evaluated at new x.

3.2 Loss functions and their pseudo-residuals

Loss	$\Psi(y, f)$	$-\partial\Psi/\partial f$	F_0
squared (L_2)	$\frac{1}{2}(y-f)^2$	$y - f$	mean
absolute (L_1)	$\lvert y-f\rvert$	$\operatorname{sign}(y-f)$	median
Huber	$\frac{1}{2}(y-f)^2$ if $\lvert y-f\rvert \le \delta$; $\delta(\lvert y-f\rvert - \delta/2)$ otherwise	$y-f$ or $\delta\operatorname{sign}(y-f)$	robust centre
quantile	$(1-\alpha)\lvert y-f\rvert$ if $y-f \le 0$; $\alpha\lvert y-f\rvert$ if $y-f>0$	α or $-(1-\alpha)$	α-quantile
Bernoulli	$\log(1+e^{-2\tilde{y}f})$, $\tilde{y}\in\{-1,1\}$	$\dfrac{2\tilde{y}}{1+e^{2\tilde{y}f}}$	$\frac{1}{2}\log\frac{\bar{p}}{1-\bar{p}}$
exponential (AdaBoost)	$e^{-\tilde{y}f}$	$\tilde{y}e^{-\tilde{y}f}$	$\frac{1}{2}\log\frac{\bar{p}}{1-\bar{p}}$

In the Bernoulli loss, f is half the log-odds ($p = 1/(1+e^{-2f})$); with f equal to the log-odds the same loss is the usual logistic log loss with pseudo-residual $y-p$, $y \in \{0,1\}$. With absolute or Huber loss a few large residuals cannot dominate the trees (Ch. 2).

3.3 Shrinkage, subsampling and early stopping

Small learning rates ($\eta = 0.01$–0.1) need more trees but generalise better: each tree only partially corrects the residual, which acts like regularisation along the path. **Stochastic gradient boosting** fits each tree on a random subsample of rows (and/or columns), which reduces correlation and variance. The number of trees B is chosen by **early stopping** on a validation set: stop when validation loss has not improved for a given number of rounds. Learning rate, number of trees, tree depth (interaction order), subsampling and early stopping control the *same* overfitting problem from different places.

3.4 Limitations of the standard GBM

Standard GBM predicts very well but has practical limits: with too many trees it can overfit; it contains no explicit regularisation of tree complexity; sequential tree fitting can be computationally inefficient. Modern implementations address these.

4 XGBoost: Regularised Second-Order Boosting

XGBoost improves both the *mathematical optimisation* and the *system computation*:

- **Regularisation**: GBM has no built-in complexity penalty; XGBoost adds L_1/L_2 penalties of the leaf weights and a per-leaf cost to the objective.
- **Second-order optimisation**: GBM follows only the first derivative (gradient) and can converge slowly; XGBoost uses a second-order Taylor expansion, i.e. the curvature (Hessian), to reach the optimum faster and more precisely.

Theorem: optimal leaf weights and split gain

At round t the new tree $f_t(x) = w_{q(x)}$ (leaf index $q(x)$, T leaves) minimises

$$\mathcal{L}^{(t)} = \sum_i L(y_i, \hat{y}_i^{(t-1)} + f_t(x_i)) + \gamma T + \tfrac{1}{2}\lambda \sum_{j=1}^{T} w_j^2.$$

With $g_i = \partial_{\hat{y}} L(y_i, \hat{y}_i^{(t-1)})$, $h_i = \partial^2_{\hat{y}} L(y_i, \hat{y}_i^{(t-1)})$, $G_j = \sum_{i \in I_j} g_i$, $H_j = \sum_{i \in I_j} h_i$, the second-order approximation gives

$$w_j^\star = -\frac{G_j}{H_j + \lambda}, \qquad \tilde{\mathcal{L}}^\star = -\frac{1}{2} \sum_{j=1}^{T} \frac{G_j^2}{H_j + \lambda} + \gamma T,$$

and the gain of splitting a leaf into L and R is

$$\text{Gain} = \frac{1}{2}\Big[\frac{G_L^2}{H_L + \lambda} + \frac{G_R^2}{H_R + \lambda} - \frac{(G_L + G_R)^2}{H_L + H_R + \lambda}\Big] - \gamma.$$

Proof. Taylor: $L(y_i, \hat{y} + f) \approx L(y_i, \hat{y}) + g_i f + \frac{1}{2} h_i f^2$. Dropping constants and grouping by leaves, $\tilde{\mathcal{L}} = \sum_j \{G_j w_j + \frac{1}{2}(H_j + \lambda) w_j^2\} + \gamma T$, a separate quadratic in each w_j; its minimiser is $-G_j/(H_j + \lambda)$ with value $-\frac{1}{2} G_j^2/(H_j + \lambda)$. The gain is the objective of one leaf minus that of two leaves, minus the extra leaf cost γ. ■

Example: what $w^\star$ is for familiar losses

Squared loss $\frac{1}{2}(y - \hat{y})^2$: $g_i = \hat{y}_i - y_i$, $h_i = 1$, so $w_j^\star = \sum_{i \in I_j} (y_i - \hat{y}_i)/(n_j + \lambda)$, a *shrunken* mean residual (Ridge on leaf values). Logistic loss on the log-odds scale: $g_i = p_i - y_i$, $h_i = p_i(1 - p_i)$, so $w_j^\star$ is one Newton step for the leaf. The constraint `min_child_weight` bounds H_j from below; γ prunes splits whose gain is too small.

On the system side: sparsity-aware split finding learns a default direction for missing values, approximate (quantile-sketch) split candidates, column blocks and cache-aware parallel computation.

5 LightGBM

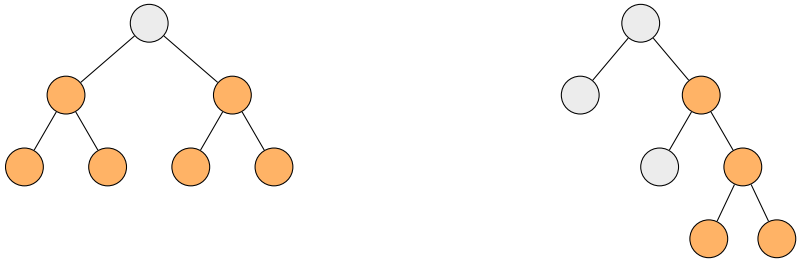
LightGBM is a highly efficient gradient-boosting implementation optimised for large and high-dimensional data:

- **Histogram-based learning**: continuous variables are bucketed into a small number of bins, drastically reducing computation and memory.
- **Leaf-wise tree growth**: the tree is not grown symmetrically; only the leaf whose split reduces loss most is expanded, deeply and selectively.
- **Efficient computation**: combining the two cuts training time and memory dramatically.

Aspect	Standard GBM	LightGBM
split search	pre-sorts every continuous value and scans all of them; bottleneck on large data	integer bins (histogram) cut computation and memory
tree growth	level-wise: grows all nodes of a level symmetrically, even unnecessary ones	leaf-wise: expands only the node that reduces the error most

Derivation: why histograms are cheap

With k bins (default 255), building a node's histogram of (G, H) sums costs $O(\#\text{rows in node})$ per feature, and finding the best split costs only $O(k)$ instead of $O(\#\text{rows})$. Moreover a child's histogram is the parent's minus the sibling's (*histogram subtraction*), so only the smaller child needs to be built. Coarse bins act as mild regularisation.



level-wise (GBM, XGBoost default) leaf-wise (LightGBM)

Figure 1: Orange nodes are created by the most recent splits. Level-wise growth splits every node of a depth before going deeper; leaf-wise growth always splits the single leaf with the largest loss reduction, giving deep, unbalanced trees for the same number of leaves.

Remark: leaf-wise growth and overfitting

For the same number of leaves, leaf-wise growth attains lower training loss than level-wise growth, but it produces deep, unbalanced trees that overfit small data; control `num_leaves`, `max_depth`, `min_data_in_leaf`. LightGBM also uses gradient-based one-side sampling (GOSS: keep rows with large $|g_i|$, subsample the rest with reweighting) and exclusive feature bundling (EFB: merge sparse features that are rarely nonzero simultaneously).

6 CatBoost

6.1 The curse of dimensionality of categorical variables

- GBM: hundreds of categories require one-hot encoding into hundreds of dummy columns; memory and trees become inefficient.
- CatBoost: however many categories, each is replaced *within a single column* by the proportion of target $= 1$ (target statistics), preventing the explosion of dimensions.

Example: one-hot (GBM) versus target statistics (CatBoost)

Category	Share of target $= 1$ in that category	CatBoost value (one column)
A	80 of 100 (80%)	$\to 0.80$
B	10 of 40 (25%)	$\to 0.25$
C	100 of 200 (50%)	$\to 0.50$
D	5 of 50 (10%)	$\to 0.10$

Instead of four new columns for A, B, C, D, one continuous column $0.80, 0.25, \ldots$ replaces them.

6.2 Naive target statistics leak the target

Theorem: target leakage of in-sample target encoding

If each row's encoding uses its own target, $\hat{x}_i = \frac{1}{n_{c_i}} \sum_{j:c_j=c_i} y_j$, then for a category observed once the encoded value equals y_i. In the extreme of all categories unique (e.g. an ID), the encoded feature equals the target on the training data and a single split fits it perfectly, while on new data the feature carries no information.

More generally, $\mathbb{E}[\hat{x}_i \mid y_i]$ depends on y_i through the term y_i/n_{c_i}, so the distribution of the encoded feature differs between training rows (which saw their own target) and test rows (*target shift*). CatBoost fixes this in two steps:

- **smoothing with a prior**: $\hat{x}_i = \dfrac{\sum_{j\in\mathcal{P}_i} \mathbb{I}(c_j = c_i)\, y_j + a\,\bar{p}}{\sum_{j\in\mathcal{P}_i} \mathbb{I}(c_j = c_i) + a}$, shrinking rare categories toward the global rate $\bar{p}$;
- **ordered target statistics**: draw a random permutation σ and let $\mathcal{P}_i = \{j : \sigma(j) < \sigma(i)\}$, i.e. each row is encoded using only rows that come *before* it, as in an online stream. A row's own target never enters its encoding, mimicking the test situation.

The same idea, **ordered boosting**, computes the residuals of a row with a model that has not seen that row, removing the "prediction shift" of standard GBM. CatBoost also uses symmetric (oblivious) trees, the same split at every node of a level, which are fast to evaluate and regularise the trees.

7 Summary and Practical Rules

Algorithm	Direction of development
XGBoost	explicit regularisation controls overfitting; second-order information (Hessian) gives precise optimisation
LightGBM	histograms and leaf-wise growth drastically reduce computation on large numerical/sparse data
CatBoost	ordered target statistics handle leakage and instability of categorical features

Rules for using boosting.

- Boosting adds shallow trees sequentially, updating the predictor in the direction that most reduces the current loss.
- Learning rate, number of iterations, tree complexity, subsampling and early stopping control the same overfitting problem from different places.
- Choose a model not by its name or one best score but by the performance and error structure repeated under a fixed validation procedure (time and group structure of Ch. 1 still apply; early stopping uses a validation fold, not the test set).