Chapter 5. Potential Problems in Regression

Lecture 6-2 (assumptions, non-Gaussianity, outliers and robust regression, breakdown point, measurement error)

Contents

1 Assumptions and Common Issues **2**

2 Non-Gaussianity: Normality Is Not an Issue **2**

2.1 Large-sample inference 2

3 Outliers and Robust Regression **3**

3.1 Why OLS is sensitive 3

3.2 Alternative objectives 3

3.3 Breakdown point 3

4 Measurement Error **4**

4.1 Remedies 4

1 Assumptions and Common Issues

The linear model relies on: linearity; normally distributed errors; independent errors; constant variance; no outliers or high-leverage points; small collinearity.

Outliers are observations with unusually large residuals; **high-leverage points** have unusual predictor values; **correlated errors** violate independence.

Kind of problem	Examples
data problems	outliers, measurement error, missing data
model problems	omitted-variable bias, misspecification, multicollinearity
error problems	non-normality, heteroscedasticity, autocorrelation

Remark: diagnostics in brief

Nonlinearity: residuals versus fitted values show a pattern (fix: transformations, polynomials, splines). **Heteroscedasticity**: a funnel in the residual plot (fix: transform Y, e.g. $\log Y$ or $\sqrt{Y}$, or weighted least squares). **Correlated errors**: tracking in residuals over time; standard errors become too small. **Leverage**: $h_{ii} = x_i^\top (X^\top X)^{-1} x_i$, average $(p+1)/n$; flag $h_{ii} \gg (p+1)/n$. **Collinearity**: variance inflation factor $\mathrm{VIF}(\hat\beta_j) = 1/(1 - R^2_{X_j|X_{-j}})$; values above 5–10 are problematic.

2 Non-Gaussianity: Normality Is Not an Issue

Theorem: exact distribution under Gaussian errors

If $y = X\beta + \varepsilon$ with $\varepsilon \sim N(0, \sigma^2 I)$ and X of full column rank, then

$$\hat\beta \sim N(\beta, \sigma^2 (X^\top X)^{-1}), \qquad \hat\beta_j \sim N(\beta_j, \sigma^2 v_j),$$

where v_j is the jth diagonal element of $(X^\top X)^{-1}$. Replacing σ by $\hat\sigma = \sqrt{\mathrm{RSS}/(n-p)}$ (p columns of X including the intercept),

$$t_j = \frac{\hat\beta_j - \beta_j}{\hat\sigma \sqrt{v_j}} \sim t_{n-p},$$

which gives tests and confidence intervals.

Proof. $\hat\beta = \beta + (X^\top X)^{-1} X^\top \varepsilon$ is a linear map of a Gaussian vector, with mean β and covariance $\sigma^2 (X^\top X)^{-1}$. The residual vector $(I - H)\varepsilon$ is independent of $\hat\beta$ (uncorrelated jointly Gaussian), and $\mathrm{RSS}/\sigma^2 = \varepsilon^\top (I-H)\varepsilon/\sigma^2 \sim \chi^2_{n-p}$ since $I - H$ is a projection of rank $n - p$. The ratio of a standard normal and an independent $\sqrt{\chi^2_{n-p}/(n-p)}$ is t_{n-p}. ■

2.1 Large-sample inference

If the errors are not normal, $\hat\beta$ is not exactly normal, but by the CLT

$$\hat\beta \;\dot\sim\; N(\beta, \sigma^2 (X^\top X)^{-1}) \quad \text{for large } n,$$

(provided errors are i.i.d. with finite variance and the design is not dominated by a few points, i.e. $\max_i h_{ii} \to 0$). Hence t-tests and confidence intervals remain approximately valid. Alternatively use

bootstrap inference: resample the data with replacement, recompute $\hat{\beta}$ for each bootstrap sample, and use the empirical distribution of $\hat{\beta}^*$ for standard errors and intervals (Ch. 3). Resampling pairs (x_i, y_i) is also valid under heteroscedasticity.

3 Outliers and Robust Regression

3.1 Why OLS is sensitive

OLS minimizes $\sum_{i=1}^{n}(y_i - x_i^\top \beta)^2$. Squaring residuals amplifies large errors, so a single extreme observation can strongly influence the fitted model. This motivates robust regression.

3.2 Alternative objectives

- **Least absolute deviation (LAD)**: $\hat{\beta} = \arg\min_\beta \sum_i |y_i - x_i^\top \beta|$; the conditional median (Ch. 2 analogue: absolute loss is minimized by the median).
- **Huber regression**: $\hat{\beta} = \arg\min_\beta \sum_i \rho_c(r_i)$, $r_i = y_i - x_i^\top \beta$,

$$\rho_c(r) = \begin{cases} \frac{1}{2}r^2, & |r| \le c, \\ c|r| - \frac{1}{2}c^2, & |r| > c. \end{cases}$$

- **Least trimmed squares (LTS)**: $\hat{\beta} = \arg\min_\beta \sum_{i=1}^{h} r^2_{(i)}$, using only the h smallest squared residuals ($r^2_{(1)} \le r^2_{(2)} \le \cdots$).

Derivation: Huber as iteratively reweighted least squares

ρ_c is convex with derivative $\psi_c(r) = r$ for $|r| \le c$ and $c\,\mathrm{sign}(r)$ otherwise. The estimating equation $\sum_i \psi_c(r_i)x_i = 0$ equals the weighted normal equation $\sum_i w_i r_i x_i = 0$ with $w_i = \min(1, c/|r_i|)$, so one iterates $\beta \leftarrow (X^\top W X)^{-1} X^\top W y$. Large residuals get weights $c/|r_i| < 1$: their influence is bounded by c, whereas under OLS it is unbounded. Usually $c = 1.345\hat{s}$ for a robust scale $\hat{s}$ (95% efficiency under normal errors).

3.3 Breakdown point

Definition: breakdown point

The (finite-sample) breakdown point of an estimator is the smallest fraction of contaminated observations that can make the estimate arbitrarily large (break it down).

OLS	LAD	Huber M-estimator	least median of squares	least trimmed squares
0	0	0	0.5	≈ 0.5

OLS can be destroyed by a single outlier (breakdown $1/n \to 0$): moving one y_i to infinity moves $\hat{\beta}$ to infinity. LAD and Huber resist outliers in y, but a single bad *leverage point* (outlying x_i) can still break them, hence breakdown 0. High-breakdown estimators (LMS: minimize the median squared residual; LTS with $h \approx n/2$) tolerate nearly 50% contamination — more than half bad points can never be distinguished from good ones, so 0.5 is the maximum.

4 Measurement Error

Suppose the true model is $Y = \beta_0 + \beta_1 X + \varepsilon$, but the predictor is measured with error:

$$\tilde{X} = X + u, \qquad u \text{ independent of } (X, \varepsilon),\ \mathbb{E}u = 0.$$

Regressing Y on the observed $\tilde{X}$ gives a biased estimator.

Theorem: attenuation bias

$$\hat{\beta}_1 \xrightarrow{p} \beta_1 \frac{\mathrm{Var}(X)}{\mathrm{Var}(X) + \mathrm{Var}(u)} = \lambda \beta_1, \qquad 0 < \lambda \leq 1.$$

Proof. The OLS slope converges to $\mathrm{Cov}(\tilde{X}, Y)/\mathrm{Var}(\tilde{X})$. Here $\mathrm{Cov}(X + u, \beta_0 + \beta_1 X + \varepsilon) = \beta_1 \mathrm{Var} X$ and $\mathrm{Var}(\tilde{X}) = \mathrm{Var} X + \mathrm{Var} u$. ■

The slope is shrunk toward zero (**attenuation bias**); λ is the reliability ratio. Measurement error in Y only, in contrast, inflates the error variance but does not bias the slope. In multiple regression, error in one predictor also biases the coefficients of correlated predictors.

4.1 Remedies

Method	Idea
instrumental variables	use a variable Z correlated with X but independent of the measurement error: $\hat{\beta}_1^{\mathrm{IV}} = \widehat{\mathrm{Cov}}(Z, Y)/\widehat{\mathrm{Cov}}(Z, \tilde{X})$
repeated measurements	average m noisy measurements, reducing $\mathrm{Var} u$ to $\mathrm{Var} u/m$; replicates also estimate $\mathrm{Var} u$ so that $\hat{\beta}_1/\hat{\lambda}$ corrects the bias
errors-in-variables models	model the measurement process explicitly (e.g. Deming regression)
SIMEX	add extra artificial noise of increasing size, observe how the estimate changes, and extrapolate back to zero measurement error
validation data	use a subset where the true X is observed to calibrate $\tilde{X}$ (regression calibration)

Summary

- Regression problems arise from the data (outliers, measurement error, missing data), the model (omitted variables, misspecification, collinearity) and the error structure (heteroscedasticity, autocorrelation).
- Normality is often not critical: the CLT gives large-sample inference and the bootstrap is an alternative.
- Outliers motivate robust regression (LAD, Huber, LTS); the breakdown point measures robustness, and only high-breakdown methods resist leverage points.
- Error in a predictor causes attenuation bias toward zero; IV, replicates, SIMEX or validation data correct it.