Chapter 10. Discrimination and Classification

Bayes rule, LDA/QDA, Fisher's discriminant, logistic regression, assessing classifiers

Contents

1 The Bayes Classification Rule

1.1 Setting

Let $Y \in \{1, \dots, K\}$ be the group and $\mathbf{X} \in \mathbb{R}^p$ the features, with priors $P(Y = k) = \pi_k$ and $\mathbf{X} \mid Y = k \sim f_k$. For $K = 2$: $\mathbf{X} \mid Y = 1 \sim f_1(\mathbf{x})$, $\mathbf{X} \mid Y = 2 \sim f_2(\mathbf{x})$, $P(Y = 1) = \pi_1$, and

$$\eta(\mathbf{x}) := P(Y = 1 \mid \mathbf{X} = \mathbf{x}) = \frac{f_1(\mathbf{x})\pi_1}{f_1(\mathbf{x})\pi_1 + f_2(\mathbf{x})\pi_2}.$$

The misclassification probability of a classifier δ (or ϕ) : $\mathbb{R}^p \to \{1, \dots, K\}$ is $P(\phi(\mathbf{X}) \neq Y) = \sum_k P(\phi(\mathbf{X}) \neq k \mid Y = k)P(Y = k)$.

Theorem 1: optimality of the Bayes rule

$\phi_{\text{Bayes}}(\mathbf{x}) = \arg\max_k P(Y = k \mid \mathbf{X} = \mathbf{x}) = \arg\max_k \pi_k f_k(\mathbf{x})$ minimises the misclassification probability: $P(\phi_{\text{Bayes}}(\mathbf{X}) \neq Y) \leq P(\phi(\mathbf{X}) \neq Y)$ for every ϕ.

$K = 2$.

$$\begin{aligned}
\text{Risk}(\delta) &= \mathrm{E}\big[\mathbf{1}_{\{Y \neq \delta(\mathbf{X})\}}\big] = \mathrm{E}_Y\Big[\mathrm{E}_{\mathbf{X}|Y}\big[\mathbf{1}_{\{Y \neq \delta(\mathbf{X})\}} \mid Y\big]\Big] \\
&= \pi_1 \mathrm{E}_{\mathbf{X}|Y=1}\big[\mathbf{1}_{\{\delta(\mathbf{X}) \neq 1\}} \mid Y = 1\big] + \pi_2 \mathrm{E}_{\mathbf{X}|Y=2}\big[\mathbf{1}_{\{\delta(\mathbf{X}) \neq 2\}} \mid Y = 2\big] \\
&= \int \underbrace{\Big[\pi_1 \mathbf{1}_{\{\delta(\mathbf{x}) \neq 1\}} f_1(\mathbf{x}) + \pi_2 \mathbf{1}_{\{\delta(\mathbf{x}) \neq 2\}} f_2(\mathbf{x})\Big]}_{=:h(\mathbf{x})} d\mathbf{x}.
\end{aligned}$$

At each $\mathbf{x}$, $\delta(\mathbf{x}) = 1$ gives $h(\mathbf{x}) = \pi_2 f_2(\mathbf{x})$ and $\delta(\mathbf{x}) = 2$ gives $h(\mathbf{x}) = \pi_1 f_1(\mathbf{x})$. To minimise the integral choose the smaller pointwise: if $\pi_2 f_2(\mathbf{x}) < \pi_1 f_1(\mathbf{x})$ choose $\delta(\mathbf{x}) = 1$; if $>$, choose $\delta(\mathbf{x}) = 2$. Thus $\phi_{\text{Bayes}}(\mathbf{x}) = \arg\max_k P(Y = k \mid \mathbf{X} = \mathbf{x})$. For general K the same argument minimises $1 - P(Y = \phi(\mathbf{x}) \mid \mathbf{x})$ pointwise. ■

Remark 2: misclassification costs

If $c(2 \mid 1)$ is the cost of classifying a member of group 1 into group 2, the rule minimising expected cost classifies into group 1 when $\dfrac{f_1(\mathbf{x})}{f_2(\mathbf{x})} \geq \dfrac{c(1 \mid 2)}{c(2 \mid 1)} \dfrac{\pi_2}{\pi_1}$.

2 Normal Populations: LDA and QDA

2.1 Common covariance: linear discriminant analysis

If $f_k = N_p(\boldsymbol{\mu}_k, \boldsymbol{\Sigma})$ ($\boldsymbol{\Sigma}_k = \boldsymbol{\Sigma}$), then $\log\{\pi_k f_k(\mathbf{x})\} = -\frac{1}{2}\mathbf{x}^\top \boldsymbol{\Sigma}^{-1}\mathbf{x} + \boldsymbol{\mu}_k^\top \boldsymbol{\Sigma}^{-1}\mathbf{x} - \frac{1}{2}\boldsymbol{\mu}_k^\top \boldsymbol{\Sigma}^{-1}\boldsymbol{\mu}_k + \log \pi_k +$ const, and the first term does not depend on k, so

$$\phi(\mathbf{x}) = \arg\min_k \delta_k(\mathbf{x}), \qquad \delta_k(\mathbf{x}) = b_{0k} - \mathbf{b}_k^\top \mathbf{x}, \quad b_{0k} = \tfrac{1}{2}\boldsymbol{\mu}_k^\top \boldsymbol{\Sigma}^{-1}\boldsymbol{\mu}_k - \log(\pi_k), \quad \mathbf{b}_k = \boldsymbol{\Sigma}^{-1}\boldsymbol{\mu}_k.$$

The discriminant function is linear in $\mathbf{x}$: **linear discriminant analysis (LDA)**. For $K = 2$, $\delta_1 < \delta_2$ rearranges to

$$\phi(\mathbf{x}) = 1 \;\text{ if }\; \mathbf{b}^\top\Big(\mathbf{x} - \frac{\boldsymbol{\mu}_1 + \boldsymbol{\mu}_2}{2}\Big) < \log\Big(\frac{\pi_1}{\pi_2}\Big), \qquad \mathbf{b} = \boldsymbol{\Sigma}^{-1}(\boldsymbol{\mu}_2 - \boldsymbol{\mu}_1).$$

$K = 2$. $\delta_1 < \delta_2 \iff (\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2)^\top \boldsymbol{\Sigma}^{-1}\mathbf{x} - \frac{1}{2}(\boldsymbol{\mu}_1^\top \boldsymbol{\Sigma}^{-1}\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2^\top \boldsymbol{\Sigma}^{-1}\boldsymbol{\mu}_2) > \log(\pi_2/\pi_1)$. Since $\boldsymbol{\mu}_1^\top \boldsymbol{\Sigma}^{-1}\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2^\top \boldsymbol{\Sigma}^{-1}\boldsymbol{\mu}_2 = (\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2)^\top \boldsymbol{\Sigma}^{-1}(\boldsymbol{\mu}_1 + \boldsymbol{\mu}_2)$, this is $-\mathbf{b}^\top(\mathbf{x} - \frac{\boldsymbol{\mu}_1 + \boldsymbol{\mu}_2}{2}) > \log(\pi_2/\pi_1)$. ■

The sample version plugs in $\boldsymbol{\mu}_k \to \bar{\mathbf{x}}_k$, $\boldsymbol{\Sigma} \to \mathbf{S}_p$ (pooled covariance), $\pi_k \to n_k/n$: $\phi(\mathbf{x}) = \arg\max_k(b_{0k} - \mathbf{b}_k^\top \mathbf{x})$-type rule with $\hat{\mathbf{b}}_k = \mathbf{S}_p^{-1}\bar{\mathbf{x}}_k$, and for $K = 2$, $\hat{\mathbf{b}} = \mathbf{S}_p^{-1}(\bar{\mathbf{x}}_2 - \bar{\mathbf{x}}_1)$.

2.2 Different covariances: quadratic discriminant analysis

If $f_k = N_p(\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$,

$$\phi(\mathbf{x}) = \arg\max_k \Big\{ -\tfrac{1}{2}\log|\boldsymbol{\Sigma}_k| - \tfrac{1}{2}(\mathbf{x}-\boldsymbol{\mu}_k)^\top \boldsymbol{\Sigma}_k^{-1}(\mathbf{x}-\boldsymbol{\mu}_k) + \log \pi_k \Big\} = \arg\max_k \left(b_{0k} - \mathbf{b}_k^\top \mathbf{x} + \mathbf{x}^\top \mathbf{C}_k \mathbf{x}\right)$$

with $\mathbf{C}_k = -\frac{1}{2}\boldsymbol{\Sigma}_k^{-1}$. The boundaries are quadratic surfaces: **QDA**. It needs $K \cdot p(p+1)/2$ more parameters, so with small samples LDA is often more stable (bias–variance trade-off).

3 Fisher's Linear Discriminant

Without distributional assumptions, find the projection that best separates the groups.

Definition 3: Fisher's LDA direction ($K = 2$)

$$\mathbf{u}_0 = \arg\max_{\mathbf{u}\in\mathbb{R}^p} \frac{(\mathbf{u}^\top \bar{\mathbf{x}}_1 - \mathbf{u}^\top \bar{\mathbf{x}}_2)^2}{\mathbf{u}^\top \mathbf{S}_p \mathbf{u}} :$$

squared difference of projected means over projected within-group variance.

Theorem 4: solution

$\mathbf{u}_0 \propto \mathbf{S}_p^{-1}(\bar{\mathbf{x}}_1 - \bar{\mathbf{x}}_2)$ (equivalently $\mathbf{S}_p^{-1}(\bar{\mathbf{x}}_2 - \bar{\mathbf{x}}_1)$), and the maximum is $\mathbf{d}^\top \mathbf{S}_p^{-1}\mathbf{d}$ with $\mathbf{d} = \bar{\mathbf{x}}_1 - \bar{\mathbf{x}}_2$, the squared sample Mahalanobis distance.

Proof. In general, to maximise $\frac{\mathbf{u}^\top \mathbf{A}\mathbf{u}}{\mathbf{u}^\top \mathbf{B}\mathbf{u}}$ ($\mathbf{B} \succ 0$) use $L(\mathbf{u}, \lambda) = \mathbf{u}^\top \mathbf{A}\mathbf{u} - \lambda(\mathbf{u}^\top \mathbf{B}\mathbf{u} - 1)$: $\partial L/\partial \mathbf{u} = \mathbf{0} \Rightarrow \mathbf{A}\mathbf{u} = \lambda \mathbf{B}\mathbf{u} \Rightarrow \mathbf{B}^{-1}\mathbf{A}\mathbf{u} = \lambda \mathbf{u}$, so the top eigenvector of $\mathbf{B}^{-1}\mathbf{A}$ is optimal. With $\mathbf{A} = \mathbf{d}\mathbf{d}^\top$ (rank 1) and $\mathbf{B} = \mathbf{S}_p$, $\mathbf{S}_p^{-1}\mathbf{d}\mathbf{d}^\top$ has a single nonzero eigenvalue $\mathbf{d}^\top \mathbf{S}_p^{-1}\mathbf{d}$ with unique eigenvector $\mathbf{S}_p^{-1}\mathbf{d}$ (rank 1). Directly: with $\mathbf{v} = \mathbf{S}_p^{1/2}\mathbf{u}$, Cauchy–Schwarz gives $(\mathbf{u}^\top \mathbf{d})^2/\mathbf{u}^\top \mathbf{S}_p \mathbf{u} = (\mathbf{v}^\top \mathbf{S}_p^{-1/2}\mathbf{d})^2/\mathbf{v}^\top \mathbf{v} \le \mathbf{d}^\top \mathbf{S}_p^{-1}\mathbf{d}$. ∎

Remark 5: Fisher's rule = LDA with equal priors

With the discriminant score $y_i = \mathbf{b}_Y^\top \mathbf{x}_i$, $\mathbf{b}_Y = \mathbf{S}_p^{-1}(\bar{\mathbf{x}}_2 - \bar{\mathbf{x}}_1)$, and the cut-off $b_{0Y} = \mathbf{b}_Y^\top(\bar{\mathbf{x}}_1 + \bar{\mathbf{x}}_2)/2$,

$$\text{“}\mathbf{b}_Y^\top \mathbf{y} - b_{0Y} < 0 \Rightarrow 1\text{”} \iff \text{“}\left[\mathbf{S}_p^{-1}(\bar{\mathbf{x}}_2 - \bar{\mathbf{x}}_1)\right]^\top \left(\mathbf{x} - \frac{\bar{\mathbf{x}}_1 + \bar{\mathbf{x}}_2}{2}\right) < 0 \Rightarrow 1\text{”},$$

which is sample LDA with $\pi_1 = \pi_2$. A rule derived without normality coincides with the Bayes rule under normality.

Several groups ($K \ge 3$): $\max_{\mathbf{u}} \frac{\mathbf{u}^\top \mathbf{B}\mathbf{u}}{\mathbf{u}^\top \mathbf{W}\mathbf{u}}$ ($\mathbf{B}, \mathbf{W}$ the MANOVA between/within SSP matrices) is solved by the eigenvectors of $\mathbf{W}^{-1}\mathbf{B}$; there are at most $\min(p, K-1)$ nonzero eigenvalues, the same η_i that appear in the Wilks/Pillai statistics of Ch. 6.

4 Logistic Regression

For $K = 2$, $Y \in \{0, 1\}$, model the posterior directly:

$$\operatorname{logit}(\eta(\mathbf{x})) := \log\left(\frac{\eta(\mathbf{x})}{1-\eta(\mathbf{x})}\right) = b + \boldsymbol{\beta}^\top \mathbf{x}, \qquad \eta(\mathbf{x}) = P(Y = 1 \mid \mathbf{X} = \mathbf{x}).$$

Classification: "$\eta(\mathbf{x}) > \frac{1}{2} \Rightarrow 1$" $\iff$ "$f(\mathbf{x}) = b + \boldsymbol{\beta}^\top \mathbf{x} > 0 \Rightarrow 1$" (a linear boundary). The log-likelihood is

$$\begin{aligned}\ell(b, \boldsymbol{\beta}) &= \sum_{i=1}^{n} \left[y_i \log(\eta(\mathbf{x}_i; b, \boldsymbol{\beta})) + (1 - y_i) \log\{1 - \eta(\mathbf{x}_i; b, \boldsymbol{\beta})\} \right] \\ &= \sum_{i=1}^{n} \left[y_i (b + \boldsymbol{\beta}^\top \mathbf{x}_i) - \log\{1 + \exp(b + \boldsymbol{\beta}^\top \mathbf{x}_i)\} \right],\end{aligned}$$

which is concave and maximised by Newton–Raphson (IRLS). In the equal-covariance normal model $\log \frac{\pi_1 f_1}{\pi_0 f_0}$ is exactly linear in $\mathbf{x}$, so the logistic model holds; LDA models the joint distribution, logistic regression only the conditional one.

Derivation: the second equality and the derivatives

$\log \eta = (b + \boldsymbol{\beta}^\top \mathbf{x}) - \log(1 + e^{b + \boldsymbol{\beta}^\top \mathbf{x}})$ and $\log(1 - \eta) = -\log(1 + e^{b + \boldsymbol{\beta}^\top \mathbf{x}})$, so $y \log \eta + (1 - y) \log(1 - \eta) = y(b + \boldsymbol{\beta}^\top \mathbf{x}) - \log(1 + e^{b + \boldsymbol{\beta}^\top \mathbf{x}})$. The gradient is $\sum_i (y_i - \eta_i)(1, \mathbf{x}_i^\top)^\top$ and the Hessian $-\sum_i \eta_i (1 - \eta_i)(1, \mathbf{x}_i^\top)^\top (1, \mathbf{x}_i^\top) \preceq 0$.

5 Simple Classifiers in High Dimensions

If $p > n$, $\mathbf{S}_p$ is singular and LDA cannot be used (e.g. the Golub data).

- **Nearest centroid**: LDA with $\boldsymbol{\Sigma} = \mathbf{I}$. $\phi(\mathbf{x}) = 1$ if $\mathbf{b}^\top (\mathbf{x} - \frac{\bar{\mathbf{x}}_1 + \bar{\mathbf{x}}_2}{2}) < 0$, $\mathbf{b} = \bar{\mathbf{x}}_2 - \bar{\mathbf{x}}_1$ (assign to the closer mean).
- **Naive Bayes classifier (diagonal LDA)**: assume independent variables and replace $\mathbf{S}_p$ by $\mathbf{D}_p = \operatorname{diag}(\mathbf{S}_p)$: $\phi(\mathbf{x}) = 1$ if $\mathbf{b}^\top (\mathbf{x} - \frac{\bar{\mathbf{x}}_1 + \bar{\mathbf{x}}_2}{2}) < 0$, $\mathbf{b} = \mathbf{D}_p^{-1} (\bar{\mathbf{x}}_2 - \bar{\mathbf{x}}_1)$.
- **k-nearest neighbours**: majority vote among the k training points closest to $\mathbf{x}$; non-parametric, high variance for small k and high bias for large k.

6 Assessment of Classifiers

- True error rate $P(\phi(\mathbf{X}) \neq Y) = \sum_{k=1}^{K} P(\phi(\mathbf{X}) \neq k \mid Y = k) P(Y = k)$; ϕ is optimal if $P(\phi(\mathbf{X}) \neq Y) \leq P(\varphi(\mathbf{X}) \neq Y)$ for all φ.
- Empirical misclassification rate $\hat{P}(\hat{\phi}_D(\mathbf{X}) \neq Y \mid D) = \frac{1}{M} \sum_{i=1}^{M} \mathbf{1}_{\{\hat{\phi}(\mathbf{x}_i) \neq y_i\}}$.
- **Training and test sets**: build $\hat{\phi}$ from $\mathbf{X}_{\text{tr}}$ and estimate the misclassification rate on $\mathbf{X}_{\text{test}}$. The training error (APER) is optimistically biased.
- **Confusion matrix**: actual group by predicted group (training and testing confusion matrices).
- **K-fold cross-validation**: split into K parts, use each in turn as the test set and average. $K \uparrow \Rightarrow$ bias $\downarrow$, variance $\uparrow$; $K \downarrow \Rightarrow$ the reverse. $K = n$ is LOOCV (Lachenbruch's holdout).