Lecture 4. Discrete Random Variables

Expectation, variance, binomial, Poisson, geometric, negative binomial, hypergeometric

Contents

4.1 Random variables and distribution functions **2**

4.2 Expectation **2**

4.3 Variance **3**

4.4 Bernoulli and binomial random variables **3**

4.5 The Poisson random variable **4**

4.6 Other discrete distributions **5**

4.7 Expected value of sums **5**

4.1 Random variables and distribution functions

A **random variable** is a real function on the sample space. Its **(cumulative) distribution function** is $F(x) = P(X \le x)$. A random variable is **discrete** if it takes at most countably many values $x_1, x_2, \dots$; its **probability mass function** is $p(x) = P(X = x)$, with $\sum_i p(x_i) = 1$ and $F(a) = \sum_{x_i \le a} p(x_i)$, a step function.

Proposition 4.1 (Properties of F)

F is nondecreasing, right-continuous, $F(-\infty) = 0$ and $F(\infty) = 1$. Moreover $P(a < X \le b) = F(b) - F(a)$ and $P(X < b) = \lim_{x \uparrow b} F(x)$.

Proof. Monotonicity is Proposition 2.1(b). If $x_n \downarrow b$, the events $\{X \le x_n\}$ decrease to $\{X \le b\}$; if $x_n \uparrow \infty$ they increase to S; if $x_n \uparrow b$ the events $\{X \le x_n\}$ increase to $\{X < b\}$. Apply continuity (Proposition 2.3). ∎

Example 4.1 (Coupon collecting)

Each purchase gives one of N coupon types, equally likely and independently. Let T be the number of purchases needed to collect all types. With A_j the event that type j is missing after n purchases, inclusion–exclusion gives

$$P(T > n) = P\Big(\bigcup_j A_j\Big) = \sum_{i=1}^{N-1} (-1)^{i+1} \binom{N}{i} \Big(\frac{N-i}{N}\Big)^n,$$

and $P(T = n) = P(T > n-1) - P(T > n)$.

4.2 Expectation

Definition (Expected value)

$E[X] = \sum_x x p(x)$, provided the sum converges absolutely.

It is the long-run average of repeated observations and the centre of gravity of the mass function. For the indicator I_A of an event, $E[I_A] = P(A)$.

Proposition 4.2

(a) $E[g(X)] = \sum_x g(x)p(x)$. (b) $E[aX + b] = aE[X] + b$. (c) On a countable sample space, $E[X] = \sum_{s \in S} X(s)P(\{s\})$, and hence $E[\sum_i X_i] = \sum_i E[X_i]$.

Proof. (a) Group the values x by $y = g(x)$: $\sum_y y P(g(X) = y) = \sum_y \sum_{x: g(x)=y} g(x)p(x) = \sum_x g(x)p(x)$. (b) is a special case. (c) Group outcomes s by the value $X(s)$; then linearity of the sum over s gives the second claim. ∎

In general $E[g(X)] \ne g(E[X])$.

Example 4.2 (Which question first?)

A quiz contestant may answer question 1 or 2 first, and goes on to the other only if the first is answered correctly. Question i is answered correctly with probability p_i, independently, and is worth V_i. Answering 1 first gives expected winnings $p_1(1 - p_2)V_1 + p_1 p_2 (V_1 + V_2) =$

$p_1V_1 + p_1p_2V_2$; the other order gives $p_2V_2 + p_1p_2V_1$. So question 1 should come first iff

$$\frac{p_1V_1}{1-p_1} > \frac{p_2V_2}{1-p_2}.$$

Example 4.3 (The friendship paradox)

In a network, let d_i be the number of friends of person i, and pick a random friendship (edge) and then one of its two ends at random. That person has D^* friends, and $P(D^* = k) = k\,P(D = k)/E[D]$, where D is the degree of a uniformly chosen person: a person with k friends is k times as likely to be reached. Hence

$$E[D^*] = \frac{E[D^2]}{E[D]} \geq E[D],$$

with equality only if all degrees are equal (by the variance inequality below). "Your friends have more friends than you", on average.

4.3 Variance

Definition (Variance)

With $\mu = E[X]$, $\operatorname{Var}(X) = E[(X-\mu)^2] = E[X^2] - \mu^2$. The standard deviation is $\sqrt{\operatorname{Var}(X)}$.

The second formula follows by expanding the square and using Proposition 4.2. As $\operatorname{Var}(X) \geq 0$, we get $E[X^2] \geq (E[X])^2$. Also $\operatorname{Var}(aX+b) = a^2 \operatorname{Var}(X)$. For a fair die, $E[X] = 7/2$, $E[X^2] = 91/6$, so $\operatorname{Var}(X) = 35/12$.

4.4 Bernoulli and binomial random variables

$X \sim \operatorname{Bin}(n,p)$ counts successes in n independent trials with success probability p:

$$p(k) = \binom{n}{k} p^k (1-p)^{n-k}, \qquad k = 0, \ldots, n.$$

Each specific sequence with k successes has probability $p^k(1-p)^{n-k}$ and there are $\binom{n}{k}$ such sequences. $\operatorname{Bin}(1,p)$ is the Bernoulli distribution.

Proposition 4.3

If $X \sim \operatorname{Bin}(n,p)$ then $E[X^k] = np\,E[(Y+1)^{k-1}]$ with $Y \sim \operatorname{Bin}(n-1,p)$. Hence $E[X] = np$ and $\operatorname{Var}(X) = np(1-p)$.

Proof. Using $i\binom{n}{i} = n\binom{n-1}{i-1}$,

$$E[X^k] = \sum_{i=1}^{n} i^{k-1}\, n \binom{n-1}{i-1} p^i (1-p)^{n-i} = np \sum_{j=0}^{n-1} (j+1)^{k-1} \binom{n-1}{j} p^j (1-p)^{n-1-j}.$$

$k = 1$ gives np; $k = 2$ gives $E[X^2] = np\{(n-1)p + 1\}$, so $\operatorname{Var}(X) = np(1-p)$. ∎

Since $p(k)/p(k-1) = \frac{(n-k+1)p}{k(1-p)}$, the pmf increases while $k < (n+1)p$ and then decreases: it is unimodal with mode $\lfloor (n+1)p \rfloor$. The ratio also gives a stable recursion for computing the cdf.

Example 4.4 (Chuck-a-luck)

A player bets 1 on a number; three dice are rolled and the player wins i if the number appears $i \geq 1$ times, and loses 1 otherwise. The number of appearances is $\mathrm{Bin}(3, 1/6)$, so the expected gain is

$$-\tfrac{125}{216} + 1 \cdot \tfrac{75}{216} + 2 \cdot \tfrac{15}{216} + 3 \cdot \tfrac{1}{216} = -\tfrac{17}{216} \approx -0.079.$$

Example 4.5 (k-out-of-n systems)

A system of n components, each working independently with probability p, works if at least half of them work. For which p is a 5-component system better than a 3-component one? $P_5 = \binom{5}{3}p^3q^2 + 5p^4q + p^5$ and $P_3 = 3p^2q + p^3$. One finds $P_5 - P_3 = 3p^2(p-1)^2(2p-1)$, so the larger system is better exactly when $p > \frac{1}{2}$.

4.5 The Poisson random variable

$X \sim \mathrm{Poisson}(\lambda)$ if $p(k) = e^{-\lambda}\lambda^k/k!$, $k = 0, 1, \ldots$; then $E[X] = \mathrm{Var}(X) = \lambda$ (by the same shift trick: $E[X] = \lambda \sum_j e^{-\lambda}\lambda^j/j!$ and $E[X^2] = \lambda E[X+1]$).

Proposition 4.4 (Poisson approximation)

If $n \to \infty$ and $p \to 0$ with $np = \lambda$ fixed, then $\binom{n}{k}p^k(1-p)^{n-k} \to e^{-\lambda}\lambda^k/k!$.

Proof. $\binom{n}{k}p^k(1-p)^{n-k} = \dfrac{n(n-1)\cdots(n-k+1)}{n^k} \cdot \dfrac{\lambda^k}{k!} \cdot \dfrac{(1-\lambda/n)^n}{(1-\lambda/n)^k}$, and the three factors tend to 1, $\lambda^k/k!$ and $e^{-\lambda}$. ∎

So the number of rare events among many trials (misprints on a page, people over 100, wrong numbers dialled) is approximately Poisson. The approximation remains good when the trials are only *weakly dependent* and success probabilities are small but unequal, with $\lambda = \sum p_i$ (the Poisson paradigm).

Example 4.6 (Poisson paradigm)

(a) Matching: the events "person i gets own hat" have probability $1/N$ and are weakly dependent, so the number of matches is approximately Poisson(1), as found exactly in Example 2.4. (b) Birthdays: each of the $\binom{n}{2}$ pairs shares a birthday with probability $1/365$, and pair events are weakly dependent, so $P(\text{no shared birthday}) \approx \exp\{-n(n-1)/730\}$. For $n = 23$ this is $e^{-253/365} \approx 0.500$; the exact value is 0.493.

Example 4.7 (Poisson process)

Events (earthquakes, arrivals) occur in time so that (i) counts in disjoint intervals are independent, (ii) in a short interval of length h, one event occurs with probability $\lambda h + o(h)$ and two or more with probability $o(h)$. Splitting $[0, t]$ into n pieces, the count $N(t)$ is approximately $\mathrm{Bin}(n, \lambda t/n)$, and letting $n \to \infty$,

$$P(N(t) = k) = e^{-\lambda t}\frac{(\lambda t)^k}{k!}.$$

4.6 Other discrete distributions

Geometric. The number X of trials up to and including the first success: $p(k) = (1-p)^{k-1}p$, $k \geq 1$, and $P(X > k) = (1-p)^k$. With $q = 1-p$, conditioning on the first trial gives $E[X] = p \cdot 1 + q(1 + E[X])$, so $E[X] = 1/p$; similarly $E[X^2] = p + q\,E[(1+X)^2]$ gives $E[X^2] = (2-p)/p^2$ and $\mathrm{Var}(X) = q/p^2$.

Negative binomial. The number X of trials needed for r successes: the r-th success at trial k requires $r-1$ successes in the first $k-1$ trials, so

$$p(k) = \binom{k-1}{r-1} p^r (1-p)^{k-r}, \quad k \geq r, \qquad E[X] = \frac{r}{p}, \quad \mathrm{Var}(X) = \frac{r(1-p)}{p^2}.$$

(X is a sum of r independent geometric waiting times, so Lecture 7 gives the moments at once.)

Example 4.8 (Banach match problem)

A smoker carries two boxes of N matches and each time picks a box at random. When a box is first found empty, what is the probability that the other contains exactly k matches? Call choosing the left box a "success". The left box is found empty with k matches on the right when success number $N+1$ occurs at trial $N+1+N-k$. So the probability is $2\binom{2N-k}{N}(\frac{1}{2})^{2N-k+1} = \binom{2N-k}{N}(\frac{1}{2})^{2N-k}$, the factor 2 accounting for the symmetric case.

Hypergeometric. Draw n balls without replacement from N, of which m are white; the number X of white balls drawn has

$$P(X = i) = \frac{\binom{m}{i}\binom{N-m}{n-i}}{\binom{N}{n}}, \qquad E[X] = \frac{nm}{N}, \qquad \mathrm{Var}(X) = n\frac{m}{N}\Big(1 - \frac{m}{N}\Big)\frac{N-n}{N-1}.$$

With $p = m/N$ this is the binomial mean, and the variance carries the finite-population correction $(N-n)/(N-1)$. When N is large relative to n, X is approximately $\mathrm{Bin}(n, m/N)$.

Example 4.9 (Capture–recapture)

To estimate an unknown population size N, catch and tag m animals, release them, and later catch n; suppose i are tagged. The probability of this, $P_i(N) = \binom{m}{i}\binom{N-m}{n-i}/\binom{N}{n}$, satisfies

$$\frac{P_i(N)}{P_i(N-1)} = \frac{(N-m)(N-n)}{N(N-m-n+i)} \geq 1 \iff N \leq \frac{mn}{i},$$

so the value of N that makes the observation most likely (the maximum likelihood estimate) is $\lfloor mn/i \rfloor$: the tagged fraction in the sample, i/n, is set equal to m/N.

Zeta (Zipf). $p(k) = C/k^{\alpha+1}$, $k \geq 1$, with $C = 1/\zeta(\alpha+1)$; it models heavy-tailed frequencies such as word counts.

4.7 Expected value of sums

By Proposition 4.2(c), $E[X_1 + \cdots + X_n] = E[X_1] + \cdots + E[X_n]$ for *any* random variables, independent or not. Writing a count as a sum of indicators is the main trick of the course.

Example 4.10

(a) The sum of n fair dice has mean $7n/2$. (b) If trial i succeeds with probability p_i, the expected number of successes is $\sum p_i$. (c) The hypergeometric mean: with I_j the indicator that white ball j is drawn, $E[X] = \sum_{j=1}^{m} P(I_j = 1) = m \cdot n/N$ (Example 2.1(b)).

Example 4.11 (Variance of the number of successes)

With $X = \sum_i I_i$, $X^2 = \sum_i I_i + \sum_{i \neq j} I_i I_j$, so $E[X^2] = \sum_i p_i + \sum_{i \neq j} P(I_i = I_j = 1)$. For independent trials with common p, $E[X^2] = np + n(n-1)p^2$ and $\mathrm{Var}(X) = np(1-p)$. For the hypergeometric, $P(I_i = I_j = 1) = \frac{m(m-1)}{N(N-1)}$ for two specific draws $i \neq j$, which gives the variance stated above.