Probability: Concepts and Applications

Lecture notes, Lectures 1–8

Contents

Lecture 1. Combinatorial Analysis 4

1.1 The basic principle of counting 4

1.2 Permutations 4

1.3 Combinations 5

1.4 Multinomial coefficients 6

1.5 Integer solutions of equations 6

Lecture 2. Axioms of Probability 8

2.1 Sample space and events 8

2.2 Axioms of probability 8

2.3 Equally likely outcomes 9

2.4 Probability as a continuous set function 10

Lecture 3. Conditional Probability and Independence 11

3.1 Conditional probability 11

3.2 Total probability and Bayes's formula 11

3.3 Independent events 13

3.4 $P(\cdot \mid F)$ is a probability 14

Lecture 4. Discrete Random Variables 15

4.1 Random variables and distribution functions 15

4.2 Expectation 15

4.3 Variance 16

4.4 Bernoulli and binomial random variables 16

4.5 The Poisson random variable 17

4.6 Other discrete distributions 18
4.7 Expected value of sums 19

Lecture 5. Continuous Random Variables 20
5.1 Densities 20
5.2 Expectation and variance 20
5.3 The uniform distribution 21
5.4 Normal random variables 21
5.5 Exponential random variables 22
5.6 Other continuous distributions 23
5.7 The distribution of a function of a random variable 23

Lecture 6. Jointly Distributed Random Variables 25
6.1 Joint distribution functions 25
6.2 Independent random variables 25
6.3 Sums of independent random variables 26
6.4 Conditional distributions 27
6.5 Order statistics 28
6.6 Functions of random variables: the Jacobian method 28
6.7 Exchangeable random variables 29

Lecture 7. Properties of Expectation 30
7.1 Expectation of sums 30
7.2 Covariance and variance of sums 31
7.3 Conditional expectation 31
7.4 Conditional expectation and prediction 33
7.5 Moment generating functions 33
7.6 The multivariate normal distribution 34
7.7 General definition of expectation 35

Lecture 8. Limit Theorems **36**

8.1 Markov and Chebyshev inequalities; the weak law **36**

8.2 The central limit theorem **36**

8.3 The strong law of large numbers **37**

8.4 Other inequalities **37**

8.5 The Poisson approximation, quantified **38**

Lecture 1. Combinatorial Analysis

Many probability questions on finite sample spaces with equally likely outcomes reduce to counting. This lecture collects the counting tools used throughout the course.

1.1 The basic principle of counting

Theorem 1.1 (Generalized counting principle)

If r experiments are performed in sequence so that the first has n_1 possible outcomes and, for each outcome of the first $i-1$ experiments, the i-th has n_i possible outcomes, then the r experiments together have $n_1 n_2 \cdots n_r$ possible outcomes.

Proof. For $r = 2$ list the outcomes as pairs (i, j), $1 \le i \le n_1$, $1 \le j \le n_2$: an $n_1 \times n_2$ array. Induct on r, treating the first $r-1$ experiments as one. ∎

The number n_i may not depend on earlier outcomes, but the *set* of outcomes may.

Example 1.1 (License plates)

A plate has 3 letters followed by 4 digits. There are $26^3 \cdot 10^4 = 175{,}760{,}000$ plates. If no letter and no digit may repeat, there are $26 \cdot 25 \cdot 24 \cdot 10 \cdot 9 \cdot 8 \cdot 7 = 78{,}624{,}000$. The number of functions from an n-point set to $\{0, 1\}$ is 2^n, by the same principle.

1.2 Permutations

An ordered arrangement of n distinct objects is a **permutation**; there are $n! = n(n-1)\cdots 1$ of them, with $0! = 1$. The number of ordered selections of r objects out of n is $n(n-1)\cdots(n-r+1) = n!/(n-r)!$.

If the n objects fall into r groups of $n_1, \ldots, n_r$ mutually indistinguishable objects, each distinguishable arrangement corresponds to exactly $n_1! \cdots n_r!$ permutations of labelled objects (permute within each group), so there are

$$\frac{n!}{n_1!\, n_2! \cdots n_r!}$$

distinguishable arrangements.

Example 1.2

(a) The letters of PEPPER (three P, two E, one R) give $6!/(3!\,2!\,1!) = 60$ arrangements.
(b) Four mathematics, three chemistry, two history and one language book are shelved so that books on the same subject stand together. Order the four subject blocks (4! ways), then order within each block: $4!\,(4!\,3!\,2!\,1!) = 24 \cdot 288 = 6912$ arrangements.
(c) A class of 6 men and 4 women is ranked by an exam with no ties: $10! = 3{,}628{,}800$ rankings; if men are ranked only among men and women among women, $6!\,4! = 17{,}280$.

1.3 Combinations

The number of *unordered* groups of r objects chosen from n is

$$\binom{n}{r} = \frac{n(n-1)\cdots(n-r+1)}{r!} = \frac{n!}{(n-r)!\,r!}, \qquad 0 \le r \le n,$$

since each group is counted $r!$ times among the ordered selections. We set $\binom{n}{r} = 0$ for $r < 0$ or $r > n$. Clearly $\binom{n}{r} = \binom{n}{n-r}$.

Proposition 1.2 (Pascal's identity)

$$\binom{n}{r} = \binom{n-1}{r-1} + \binom{n-1}{r} \text{ for } 1 \le r \le n.$$

Proof. Fix one object. The r-subsets containing it number $\binom{n-1}{r-1}$; those not containing it number $\binom{n-1}{r}$. ∎

Theorem 1.3 (Binomial theorem)

$$(x+y)^n = \sum_{k=0}^{n} \binom{n}{k} x^k y^{n-k}.$$

Proof. Expand $(x_1+y_1)\cdots(x_n+y_n)$: it is the sum of 2^n products, one for each choice of x_i or y_i from every factor. The products with exactly k of the x's are indexed by the k-subsets of $\{1,\dots,n\}$; there are $\binom{n}{k}$ of them. Now set $x_i = x$, $y_i = y$. ∎

Taking $x = y = 1$ shows that an n-set has $\sum_k \binom{n}{k} = 2^n$ subsets; $x = -1, y = 1$ gives $\sum_k (-1)^k \binom{n}{k} = 0$ for $n \ge 1$.

Combinatorial arguments ("count one set in two ways") prove many identities quickly.

Proposition 1.4

$$\text{(a)} \sum_{k=1}^{n} k\binom{n}{k} = n2^{n-1}, \qquad \text{(b)} \binom{n+m}{r} = \sum_{k=0}^{r} \binom{n}{k}\binom{m}{r-k} \quad \text{(Vandermonde)}.$$

Proof. (a) Count pairs (committee, chair) from n people. Choosing a committee of size k and then its chair gives the left side; choosing the chair first (n ways) and then any subset of the others (2^{n-1}) gives the right. (b) Choose r people from n men and m women by first fixing the number k of men. ∎

Example 1.3 (Antennas in a row)

Of n antennas in a line, m are defective and indistinguishable, as are the $n-m$ working ones. How many orderings have no two defectives adjacent? Line up the $n-m$ working antennas; this creates $n-m+1$ gaps (including both ends), and at most one defective may go into each gap. So the answer is $\binom{n-m+1}{m}$. For instance, $n = 4, m = 2$ gives $\binom{3}{2} = 3$: DWDW, DWWD, WDWD.

1.4 Multinomial coefficients

The number of ways to divide n distinct objects into r *labelled* groups of sizes $n_1, \dots, n_r$ ($\sum n_i = n$) is

$$\binom{n}{n_1, \dots, n_r} = \binom{n}{n_1}\binom{n-n_1}{n_2}\cdots\binom{n-n_1-\cdots-n_{r-1}}{n_r} = \frac{n!}{n_1!\cdots n_r!},$$

which is the same number as in Section 1.2: assigning object j a group label is the same as arranging a word with n_i copies of label i.

Theorem 1.5 (Multinomial theorem)

$(x_1 + \cdots + x_r)^n = \sum_{n_1+\cdots+n_r=n} \binom{n}{n_1, \dots, n_r} x_1^{n_1}\cdots x_r^{n_r}$, the sum over nonnegative integer vectors.

Example 1.4 (Labelled versus unlabelled groups)

Ten children are split into an A team and a B team of five each: $\binom{10}{5,5} = 252$ ways. If the two teams are not labelled (just "two teams of five"), each division is counted twice, giving 126.

Example 1.5 (Knockout tournament)

In a knockout tournament with $n = 2^m$ players, the first round pairs the players and each pair produces one winner. The number of pairings is $n!/\{2^{n/2}(n/2)!\}$ (arrange in a row, pair neighbours, then forget the order within and among pairs), and each pairing has $2^{n/2}$ outcomes, so the first round has $n!/(n/2)!$ possible outcomes. Repeating round by round, the whole tournament has

$$\frac{n!}{(n/2)!}\cdot\frac{(n/2)!}{(n/4)!}\cdots\frac{2!}{1!} = n!$$

outcomes.

1.5 Integer solutions of equations

Proposition 1.6 (Stars and bars)

There are $\binom{n-1}{r-1}$ vectors of *positive* integers with $x_1 + \cdots + x_r = n$, and $\binom{n+r-1}{r-1}$ vectors of *nonnegative* integers.

Proof. Write n stars in a row; a positive solution corresponds to choosing $r-1$ of the $n-1$ gaps between stars as dividers. For nonnegative solutions put $y_i = x_i + 1$: positive solutions of $\sum y_i = n + r$, of which there are $\binom{n+r-1}{r-1}$. ∎

Equivalently, $\binom{n+r-1}{r-1}$ counts the ways to put n indistinguishable balls into r distinguishable urns, and it is the number of terms in the multinomial expansion of $(x_1 + \cdots + x_r)^n$.

Example 1.6 (Investment)

\$20,000 is invested in units of \$1000 among 4 options. If all of it must be invested, there are $\binom{23}{3} = 1771$ strategies. If not all need be invested, add a fifth "cash" variable: $\binom{24}{4} = 10{,}626$. If each option must receive at least one unit, $\binom{19}{3} = 969$.

Example 1.7 (Antennas revisited)

In Example 1.3, an arrangement is determined by the numbers $x_1, \ldots, x_{m+1}$ of working antennas before the first defective, between consecutive defectives, and after the last. These satisfy $\sum x_i = n - m$ with $x_1, x_{m+1} \geq 0$ and $x_2, \ldots, x_m \geq 1$. Substituting $y_1 = x_1 + 1$, $y_{m+1} = x_{m+1} + 1$ gives a positive solution of $\sum y_i = n - m + 2$, so the count is $\binom{n-m+1}{m}$ again. Requiring at least two working antennas between defectives is now easy: substitute $y_i = x_i - 1$ for the interior terms.

Lecture 2. Axioms of Probability

2.1 Sample space and events

The **sample space** S is the set of all possible outcomes of an experiment; an **event** is a subset of S. For events E, F we use the union $E \cup F$, the intersection $EF = E \cap F$, the complement $E^c = S \setminus E$, and the empty event $\emptyset$. E and F are **mutually exclusive** if $EF = \emptyset$. Union and intersection are commutative, associative and distribute over each other, and

$$\Big(\bigcup_{i=1}^{n} E_i\Big)^c = \bigcap_{i=1}^{n} E_i^c, \qquad \Big(\bigcap_{i=1}^{n} E_i\Big)^c = \bigcup_{i=1}^{n} E_i^c \qquad \text{(De Morgan)}.$$

For the first, an outcome lies outside every E_i exactly when it lies outside their union; the second follows by applying the first to the E_i^c.

2.2 Axioms of probability

The relative-frequency idea, $P(E) = \lim n(E)/n$ over repeated trials, motivates but does not define probability; instead we impose axioms.

Definition (Probability)

A function P on events is a **probability** if (i) $0 \le P(E) \le 1$; (ii) $P(S) = 1$; (iii) for mutually exclusive $E_1, E_2, \ldots$, $P\Big(\bigcup_{i=1}^{\infty} E_i\Big) = \sum_{i=1}^{\infty} P(E_i)$.

Taking $E_1 = S$ and $E_i = \emptyset$ for $i \ge 2$ forces $P(\emptyset) = 0$, and then (iii) gives finite additivity. (For uncountable S one can only define P on a suitable class of events; we ignore this.)

Proposition 2.1

(a) $P(E^c) = 1 - P(E)$. (b) $E \subset F \Rightarrow P(E) \le P(F)$. (c) $P(E \cup F) = P(E) + P(F) - P(EF)$.

Proof. (a) E and E^c are disjoint with union S. (b) $F = E \cup E^c F$, a disjoint union. (c) $E \cup F = E \cup E^c F$ and $F = EF \cup E^c F$, both disjoint; subtract. ∎

Theorem 2.2 (Inclusion–exclusion)

$$P\Big(\bigcup_{i=1}^{n} E_i\Big) = \sum_i P(E_i) - \sum_{i<j} P(E_i E_j) + \ldots + (-1)^{r+1} \sum_{i_1 < \cdots < i_r} P(E_{i_1} \cdots E_{i_r}) + \ldots + (-1)^{n+1} P(E_1 \cdots E_n).$$

Proof. For a finite or countable S it suffices to check that each outcome contributes its probability equally to both sides. An outcome in none of the E_i contributes 0 to both. An outcome in exactly $m \ge 1$ of them contributes once on the left, and on the right $\binom{m}{1} - \binom{m}{2} + \cdots \pm \binom{m}{m} = 1 - \sum_{i=0}^{m} (-1)^i \binom{m}{i} = 1$ times. (For general S, induct on n using Proposition 2.1(c).) ∎

Truncating the sum gives alternately upper and lower bounds (Bonferroni): $P(\bigcup E_i) \le \sum P(E_i)$ and $P(\bigcup E_i) \ge \sum P(E_i) - \sum_{i<j} P(E_i E_j)$, and so on.

2.3 Equally likely outcomes

If $S = \{1, \dots, N\}$ and every outcome has probability $1/N$, then $P(E) = |E|/N$, and probability reduces to Lecture 1. It is often convenient to treat objects as distinguishable even when they are not; this never changes the answer and makes outcomes equally likely.

Example 2.1

(a) Three balls are drawn from an urn with 6 white and 5 black balls. Then $P(\text{one white, two black}) = \binom{6}{1}\binom{5}{2}/\binom{11}{3} = 60/165 = 4/11$. Counting ordered draws gives $3 \cdot 6 \cdot 5 \cdot 4/(11 \cdot 10 \cdot 9) = 4/11$ as well.
(b) k of n balls are drawn, one of them special. By symmetry, the special ball is equally likely to be in any position of a random ordering, so $P(\text{special drawn}) = k/n$; directly, $\binom{1}{1}\binom{n-1}{k-1}/\binom{n}{k} = k/n$.
(c) A poker full house: choose the rank of the triple and its 3 suits, then the rank of the pair and its 2 suits: $13\binom{4}{3} \cdot 12\binom{4}{2}/\binom{52}{5} = 3744/2{,}598{,}960 \approx 0.00144$.

Example 2.2 (Birthday problem)

Ignoring leap years and assuming all 365^n birthday assignments equally likely, the probability that n people have distinct birthdays is $365 \cdot 364 \cdots (366 - n)/365^n$. This drops below $1/2$ at $n = 23$: the probability of a shared birthday is 0.507 for $n = 23$ and 0.970 for $n = 50$. The surprise disappears when one counts pairs: 23 people form $\binom{23}{2} = 253$ pairs, each matching with probability $1/365$.

Example 2.3 (The card after the first ace)

A shuffled deck is turned up card by card. Is the card following the first ace more likely to be the ace of spades or the two of clubs? Count orderings in which the ace of spades follows the first ace: remove it, order the other 51 cards (51! ways), and reinsert it right after the first ace. So the probability is $51!/52! = 1/52$. The same argument applies to the two of clubs, so both have probability $1/52$.

Example 2.4 (The matching problem)

N people throw their hats into a pile and each picks one at random. Let E_i be the event that person i gets their own hat. For $i_1 < \cdots < i_r$, $P(E_{i_1} \cdots E_{i_r}) = (N - r)!/N!$, and there are $\binom{N}{r}$ such terms, so by inclusion–exclusion

$$P\Big(\bigcup_i E_i\Big) = \sum_{r=1}^{N} (-1)^{r+1} \binom{N}{r} \frac{(N-r)!}{N!} = \sum_{r=1}^{N} \frac{(-1)^{r+1}}{r!},$$

so that

$$P(\text{no match}) = \sum_{r=0}^{N} \frac{(-1)^r}{r!} \xrightarrow[N\to\infty]{} e^{-1} \approx 0.368.$$

Exactly k matches: choose the k people who match and require no match among the other $N - k$; this gives

$$P(\text{exactly } k) = \binom{N}{k} \frac{(N-k)!}{N!} \sum_{r=0}^{N-k} \frac{(-1)^r}{r!} = \frac{1}{k!} \sum_{r=0}^{N-k} \frac{(-1)^r}{r!} \approx \frac{e^{-1}}{k!},$$

a Poisson(1) probability (Lecture 4).

Example 2.5 (Couples at a round table)

Ten couples sit at random around a round table (19! equally likely circular orders). Let E_i be the event that couple i sit together. Gluing r given couples into blocks leaves $20-r$ units, which can be arranged in $(19-r)!$ circular orders, and each block in 2 ways. Hence

$$P(\text{no couple together}) = \sum_{r=0}^{10} (-1)^r \binom{10}{r} \frac{2^r (19-r)!}{19!} \approx 0.3395.$$

2.4 Probability as a continuous set function

A sequence of events is **increasing** if $E_n \subset E_{n+1}$, with limit $\bigcup E_n$, and **decreasing** if $E_n \supset E_{n+1}$, with limit $\bigcap E_n$.

Proposition 2.3 (Continuity)

For an increasing or decreasing sequence, $\lim_n P(E_n) = P(\lim_n E_n)$.

Proof. Increasing case: put $F_1 = E_1$ and $F_n = E_n E_{n-1}^c$. The F_n are disjoint, $\bigcup_1^n F_i = E_n$ and $\bigcup_1^\infty F_i = \bigcup_1^\infty E_i$, so by axiom (iii)

$$P\Big(\bigcup_i E_i\Big) = \sum_{i=1}^{\infty} P(F_i) = \lim_n \sum_{i=1}^{n} P(F_i) = \lim_n P(E_n).$$

Decreasing case: apply this to the increasing E_n^c and use De Morgan. ■

Example 2.6

A fair coin is tossed forever. Let E_n be the event that the first n tosses are all tails; then E_n decreases to the event T that no head ever appears, so $P(T) = \lim 2^{-n} = 0$. Thus a head appears eventually with probability 1, although "all tails" is a possible outcome: probability zero does not mean impossible.

Subjective probability (a person's degree of belief) obeys the same axioms if beliefs are to be coherent, so the theory applies to it as well.

Lecture 3. Conditional Probability and Independence

3.1 Conditional probability

Definition (Conditional probability)

If $P(F) > 0$, the conditional probability of E given F is $P(E \mid F) = P(EF)/P(F)$.

With equally likely outcomes, conditioning on F makes the outcomes in F equally likely, so $P(E \mid F) = |EF|/|F|$: we simply shrink the sample space to F.

Proposition 3.1 (Multiplication rule)

$$P(E_1E_2\cdots E_n) = P(E_1)P(E_2 \mid E_1)P(E_3 \mid E_1E_2)\cdots P(E_n \mid E_1\cdots E_{n-1}).$$

Proof. Write each factor as a ratio; the product telescopes. ∎

Example 3.1

(a) An urn holds 8 red and 4 white balls; two are drawn without replacement. $P(\text{both red}) = \frac{8}{12}\cdot\frac{7}{11} = \frac{14}{33}$, which agrees with $\binom{8}{2}/\binom{12}{2}$.
(b) A deck is divided at random into four piles of 13. Let E_i be the event that the i-th pile has exactly one ace (and put the aces in piles in order). Placing the aces one at a time among the remaining slots,

$$P(E_1E_2E_3E_4) = 1\cdot\frac{39}{51}\cdot\frac{26}{50}\cdot\frac{13}{49} \approx 0.105,$$

since given that the first ace is in some pile, the second ace must be in one of the 39 of the remaining 51 positions that lie in other piles, and so on.

Example 3.2 (Matching, recursively)

In the matching problem (Example 2.4), let P_N be the probability of no match among N people and M the event that person 1 takes their own hat. Then $P_N = P(\text{no match} \mid M^c)\frac{N-1}{N}$. Given M^c, say person 1 took hat $j \neq 1$. The other $N-1$ people choose among $N-1$ hats that include hat 1 but not hat j. Either person j takes hat 1 (probability $\frac{1}{N-1}$) and the remaining $N-2$ must have no match, or we may regard hat 1 as person j's "own" hat, and then we need no match among $N-1$ people. So

$$P_N = \frac{N-1}{N}\Big(P_{N-1} + \frac{P_{N-2}}{N-1}\Big), \qquad \text{i.e.} \qquad P_N - P_{N-1} = -\frac{1}{N}(P_{N-1} - P_{N-2}).$$

With $P_1 = 0$, $P_2 = \frac{1}{2}$, induction gives $P_N - P_{N-1} = (-1)^N/N!$ and hence $P_N = \sum_{r=0}^{N}(-1)^r/r!$ once more.

3.2 Total probability and Bayes's formula

Since $E = EF \cup EF^c$,

$$P(E) = P(E \mid F)P(F) + P(E \mid F^c)(1 - P(F)),$$

a weighted average of the conditional probabilities. More generally:

Theorem 3.2 (Total probability and Bayes's formula)

If $F_1, \dots, F_n$ are mutually exclusive with union S, then

$$P(E) = \sum_{i=1}^{n} P(E \mid F_i)P(F_i), \qquad P(F_j \mid E) = \frac{P(E \mid F_j)P(F_j)}{\sum_i P(E \mid F_i)P(F_i)}.$$

The $P(F_j)$ are **prior** probabilities of the hypotheses, and $P(F_j \mid E)$ are **posterior** probabilities after the evidence E. In odds form, for a hypothesis H,

$$\frac{P(H \mid E)}{P(H^c \mid E)} = \frac{P(H)}{P(H^c)} \cdot \frac{P(E \mid H)}{P(E \mid H^c)} :$$

new evidence multiplies the odds by the **likelihood ratio**.

Example 3.3 (A screening test)

A test detects a disease with probability 0.95 when it is present, but gives a false positive for 1% of healthy people. If 0.5% of the population has the disease,

$$P(D \mid +) = \frac{0.95 \cdot 0.005}{0.95 \cdot 0.005 + 0.01 \cdot 0.995} \approx 0.323.$$

Only a third of positives are diseased: the healthy are so numerous that their 1% of false positives outnumbers the true positives.

Example 3.4 (Multiple choice)

A student knows the answer with probability p and otherwise guesses among m alternatives. Then

$$P(\text{knew} \mid \text{correct}) = \frac{p}{p + (1-p)/m} = \frac{mp}{1 + (m-1)p}.$$

With $m = 5$, $p = 1/2$ this is $5/6$.

Example 3.5 (Three cards)

One card is red on both sides, one black on both sides, one red on one side and black on the other. A card is drawn at random and laid down; the upper side is red. The tempting answer $1/2$ for "the other side is red" is wrong: there are three equally likely red faces that could be on top, and two of them belong to the red–red card. Formally $P(RR \mid R \text{ up}) = \frac{1/3}{1/3 \cdot 1 + 1/3 \cdot 1/2} = \frac{2}{3}$.

Example 3.6 (Searching for a plane)

A plane is equally likely to be in one of three regions. If it is in region i, a search there misses it with probability β_i. After an unsuccessful search of region 1,

$$P(R_1 \mid \text{miss}) = \frac{\beta_1/3}{\beta_1/3 + 1/3 + 1/3} = \frac{\beta_1}{\beta_1 + 2}, \qquad P(R_j \mid \text{miss}) = \frac{1}{\beta_1 + 2} \quad (j = 2, 3).$$

The probability of region 1 decreases and the others increase, as intuition says.

For a fixed F with $P(F) > 0$, $Q(E) = P(E \mid F)$ is itself a probability, so all earlier results apply to it; e.g. $P(E_1 \mid F) = P(E_1 \mid E_2 F)P(E_2 \mid F) + P(E_1 \mid E_2^c F)P(E_2^c \mid F)$.

3.3 Independent events

Definition (Independence)

E and F are **independent** if $P(EF) = P(E)P(F)$. Events $E_1, \dots, E_n$ are (mutually) independent if $P(E_{i_1} \cdots E_{i_r}) = P(E_{i_1}) \cdots P(E_{i_r})$ for every subcollection.

If E and F are independent then so are E and F^c, since $P(EF^c) = P(E) - P(E)P(F) = P(E)P(F^c)$.

Example 3.7 (Pairwise does not imply mutual)

Roll two fair dice. Let E: the sum is 7; F: the first die is 4; G: the second die is 3. Each pair is independent ($P(EF) = 1/36 = P(E)P(F)$ etc.), but $P(E \mid FG) = 1 \neq P(E)$. Note also that F_6 (sum 6) is *not* independent of F while E is: $P(\text{sum} = 7 \mid \text{first} = k) = 1/6$ for every k.

Example 3.8 (Reliability)

A parallel system works if at least one of n independent components works; component i works with probability p_i. Then $P(\text{works}) = 1 - \prod_i (1 - p_i)$. A series system works with probability $\prod_i p_i$.

Example 3.9 (E before F)

Independent trials are repeated; on each, E occurs with probability $P(E)$ and F with $P(F)$, $EF = \emptyset$. Conditioning on the first trial, $P(E \text{ first}) = P(E) + (1 - P(E) - P(F))P(E \text{ first})$, so

$$P(E \text{ occurs before } F) = \frac{P(E)}{P(E) + P(F)}.$$

In craps, a first roll of 7 or 11 wins, 2, 3 or 12 loses, and any other sum k becomes the "point", after which one rolls until k (win) or 7 (lose). With $p_k = P(\text{sum} = k)$,

$$P(\text{win}) = p_7 + p_{11} + \sum_{k \in \{4,5,6,8,9,10\}} p_k \frac{p_k}{p_k + p_7} = \frac{244}{495} \approx 0.493.$$

Example 3.10 (The problem of the points)

Trials succeed independently with probability p. What is the probability of n successes before m failures? This happens iff there are at least n successes among the first $m + n - 1$ trials (either way the contest is decided by then), so

$$P_{n,m} = \sum_{k=n}^{m+n-1} \binom{m+n-1}{k} p^k (1-p)^{m+n-1-k}.$$

This was Pascal and Fermat's question of dividing the stakes of an interrupted game: the stakes should be split in the ratio $P_{n,m} : 1 - P_{n,m}$. For instance with $p = 1/2$, $n = 4$, $m = 7$ it is 0.828.

Example 3.11 (Gambler's ruin)

A gambler with fortune i wins or loses one unit on each independent play, with probabilities p and $q = 1-p$, until reaching N or 0. Let P_i be the probability of reaching N. Conditioning on the first play, $P_i = pP_{i+1} + qP_{i-1}$, i.e. $P_{i+1} - P_i = \frac{q}{p}(P_i - P_{i-1})$, with $P_0 = 0$. Hence $P_{i+1} - P_i = (q/p)^i P_1$, and summing,

$$P_i = \begin{cases} \dfrac{1-(q/p)^i}{1-(q/p)^N}, & p \neq \frac{1}{2}, \\ \dfrac{i}{N}, & p = \frac{1}{2}, \end{cases}$$

after using $P_N = 1$ to find P_1. Letting $N \to \infty$: if $p \leq \frac{1}{2}$ the gambler is ruined with probability 1 against an infinitely rich opponent; if $p > \frac{1}{2}$ survival has probability $1-(q/p)^i$.

3.4 $P(\cdot \mid F)$ is a probability

Since conditional probabilities obey the axioms, one may condition sequentially, and conditional independence makes sense: E_1, E_2 are conditionally independent given F if $P(E_1E_2 \mid F) = P(E_1 \mid F)P(E_2 \mid F)$.

Example 3.12 (Laplace's rule of succession)

A box contains $k+1$ coins; coin i shows heads with probability i/k, $i = 0, \ldots, k$. A coin is chosen at random and flipped n times, giving n heads. Given the coin, flips are independent, so

$$P(\text{head on flip } n+1 \mid n \text{ heads}) = \frac{\sum_i (i/k)^{n+1}/(k+1)}{\sum_i (i/k)^n/(k+1)} \approx \frac{\int_0^1 x^{n+1}dx}{\int_0^1 x^n dx} = \frac{n+1}{n+2}$$

for large k. The flips are not independent unconditionally: early heads are evidence that the coin is good.

Example 3.13 (Updating sequentially)

Hypotheses $H_1, \ldots, H_n$ have priors $P(H_i)$, and pieces of evidence $E_1, \ldots, E_k$ are conditionally independent given each H_i. Then

$$P(H_i \mid E_1 \cdots E_k) \propto P(H_i) \prod_{j=1}^{k} P(E_j \mid H_i),$$

so the posterior after E_1 serves as the prior for E_2, and so on; the order in which the evidence arrives does not matter.

Lecture 4. Discrete Random Variables

4.1 Random variables and distribution functions

A **random variable** is a real function on the sample space. Its **(cumulative) distribution function** is $F(x) = P(X \le x)$. A random variable is **discrete** if it takes at most countably many values $x_1, x_2, \dots$; its **probability mass function** is $p(x) = P(X = x)$, with $\sum_i p(x_i) = 1$ and $F(a) = \sum_{x_i \le a} p(x_i)$, a step function.

Proposition 4.1 (Properties of F)

F is nondecreasing, right-continuous, $F(-\infty) = 0$ and $F(\infty) = 1$. Moreover $P(a < X \le b) = F(b) - F(a)$ and $P(X < b) = \lim_{x \uparrow b} F(x)$.

Proof. Monotonicity is Proposition 2.1(b). If $x_n \downarrow b$, the events $\{X \le x_n\}$ decrease to $\{X \le b\}$; if $x_n \uparrow \infty$ they increase to S; if $x_n \uparrow b$ the events $\{X \le x_n\}$ increase to $\{X < b\}$. Apply continuity (Proposition 2.3). ∎

Example 4.1 (Coupon collecting)

Each purchase gives one of N coupon types, equally likely and independently. Let T be the number of purchases needed to collect all types. With A_j the event that type j is missing after n purchases, inclusion–exclusion gives

$$P(T > n) = P\Big(\bigcup_j A_j\Big) = \sum_{i=1}^{N-1} (-1)^{i+1} \binom{N}{i} \Big(\frac{N-i}{N}\Big)^n,$$

and $P(T = n) = P(T > n-1) - P(T > n)$.

4.2 Expectation

Definition (Expected value)

$E[X] = \sum_x x p(x)$, provided the sum converges absolutely.

It is the long-run average of repeated observations and the centre of gravity of the mass function. For the indicator I_A of an event, $E[I_A] = P(A)$.

Proposition 4.2

(a) $E[g(X)] = \sum_x g(x)p(x)$. (b) $E[aX + b] = aE[X] + b$. (c) On a countable sample space, $E[X] = \sum_{s \in S} X(s)P(\{s\})$, and hence $E[\sum_i X_i] = \sum_i E[X_i]$.

Proof. (a) Group the values x by $y = g(x)$: $\sum_y y P(g(X) = y) = \sum_y \sum_{x: g(x)=y} g(x)p(x) = \sum_x g(x)p(x)$. (b) is a special case. (c) Group outcomes s by the value $X(s)$; then linearity of the sum over s gives the second claim. ∎

In general $E[g(X)] \ne g(E[X])$.

Example 4.2 (Which question first?)

A quiz contestant may answer question 1 or 2 first, and goes on to the other only if the first is answered correctly. Question i is answered correctly with probability p_i, independently, and is worth V_i. Answering 1 first gives expected winnings $p_1(1-p_2)V_1 + p_1p_2(V_1+V_2) = p_1V_1 + p_1p_2V_2$; the other order gives $p_2V_2 + p_1p_2V_1$. So question 1 should come first iff

$$\frac{p_1V_1}{1-p_1} > \frac{p_2V_2}{1-p_2}.$$

Example 4.3 (The friendship paradox)

In a network, let d_i be the number of friends of person i, and pick a random friendship (edge) and then one of its two ends at random. That person has D^* friends, and $P(D^* = k) = k\,P(D=k)/E[D]$, where D is the degree of a uniformly chosen person: a person with k friends is k times as likely to be reached. Hence

$$E[D^*] = \frac{E[D^2]}{E[D]} \geq E[D],$$

with equality only if all degrees are equal (by the variance inequality below). "Your friends have more friends than you", on average.

4.3 Variance

Definition (Variance)

With $\mu = E[X]$, $\operatorname{Var}(X) = E[(X-\mu)^2] = E[X^2] - \mu^2$. The standard deviation is $\sqrt{\operatorname{Var}(X)}$.

The second formula follows by expanding the square and using Proposition 4.2. As $\operatorname{Var}(X) \geq 0$, we get $E[X^2] \geq (E[X])^2$. Also $\operatorname{Var}(aX+b) = a^2 \operatorname{Var}(X)$. For a fair die, $E[X] = 7/2$, $E[X^2] = 91/6$, so $\operatorname{Var}(X) = 35/12$.

4.4 Bernoulli and binomial random variables

$X \sim \operatorname{Bin}(n,p)$ counts successes in n independent trials with success probability p:

$$p(k) = \binom{n}{k} p^k (1-p)^{n-k}, \qquad k = 0, \ldots, n.$$

Each specific sequence with k successes has probability $p^k(1-p)^{n-k}$ and there are $\binom{n}{k}$ such sequences. $\operatorname{Bin}(1,p)$ is the Bernoulli distribution.

Proposition 4.3

If $X \sim \operatorname{Bin}(n,p)$ then $E[X^k] = np\,E[(Y+1)^{k-1}]$ with $Y \sim \operatorname{Bin}(n-1,p)$. Hence $E[X] = np$ and $\operatorname{Var}(X) = np(1-p)$.

Proof. Using $i\binom{n}{i} = n\binom{n-1}{i-1}$,

$$E[X^k] = \sum_{i=1}^{n} i^{k-1}\, n \binom{n-1}{i-1} p^i (1-p)^{n-i} = np \sum_{j=0}^{n-1} (j+1)^{k-1} \binom{n-1}{j} p^j (1-p)^{n-1-j}.$$

$k=1$ gives np; $k=2$ gives $E[X^2]=np\{(n-1)p+1\}$, so $\mathrm{Var}(X)=np(1-p)$. ■

Since $p(k)/p(k-1)=\frac{(n-k+1)p}{k(1-p)}$, the pmf increases while $k<(n+1)p$ and then decreases: it is unimodal with mode $\lfloor (n+1)p\rfloor$. The ratio also gives a stable recursion for computing the cdf.

Example 4.4 (Chuck-a-luck)

A player bets 1 on a number; three dice are rolled and the player wins i if the number appears $i\geq 1$ times, and loses 1 otherwise. The number of appearances is $\mathrm{Bin}(3,1/6)$, so the expected gain is

$$-\tfrac{125}{216}+1\cdot\tfrac{75}{216}+2\cdot\tfrac{15}{216}+3\cdot\tfrac{1}{216}=-\tfrac{17}{216}\approx -0.079.$$

Example 4.5 (k-out-of-n systems)

A system of n components, each working independently with probability p, works if at least half of them work. For which p is a 5-component system better than a 3-component one? $P_5=\binom{5}{3}p^3q^2+5p^4q+p^5$ and $P_3=3p^2q+p^3$. One finds $P_5-P_3=3p^2(p-1)^2(2p-1)$, so the larger system is better exactly when $p>\frac{1}{2}$.

4.5 The Poisson random variable

$X\sim\mathrm{Poisson}(\lambda)$ if $p(k)=e^{-\lambda}\lambda^k/k!$, $k=0,1,\ldots$; then $E[X]=\mathrm{Var}(X)=\lambda$ (by the same shift trick: $E[X]=\lambda\sum_j e^{-\lambda}\lambda^j/j!$ and $E[X^2]=\lambda E[X+1]$).

Proposition 4.4 (Poisson approximation)

If $n\to\infty$ and $p\to 0$ with $np=\lambda$ fixed, then $\binom{n}{k}p^k(1-p)^{n-k}\to e^{-\lambda}\lambda^k/k!$.

Proof. $\binom{n}{k}p^k(1-p)^{n-k}=\dfrac{n(n-1)\cdots(n-k+1)}{n^k}\cdot\dfrac{\lambda^k}{k!}\cdot\dfrac{(1-\lambda/n)^n}{(1-\lambda/n)^k}$, and the three factors tend to 1, $\lambda^k/k!$ and $e^{-\lambda}$. ■

So the number of rare events among many trials (misprints on a page, people over 100, wrong numbers dialled) is approximately Poisson. The approximation remains good when the trials are only *weakly dependent* and success probabilities are small but unequal, with $\lambda=\sum p_i$ (the Poisson paradigm).

Example 4.6 (Poisson paradigm)

(a) Matching: the events "person i gets own hat" have probability $1/N$ and are weakly dependent, so the number of matches is approximately Poisson(1), as found exactly in Example 2.4. (b) Birthdays: each of the $\binom{n}{2}$ pairs shares a birthday with probability $1/365$, and pair events are weakly dependent, so $P(\text{no shared birthday})\approx\exp\{-n(n-1)/730\}$. For $n=23$ this is $e^{-253/365}\approx 0.500$; the exact value is 0.493.

Example 4.7 (Poisson process)

Events (earthquakes, arrivals) occur in time so that (i) counts in disjoint intervals are independent, (ii) in a short interval of length h, one event occurs with probability $\lambda h+o(h)$ and two or more with probability $o(h)$. Splitting $[0,t]$ into n pieces, the count $N(t)$ is approximately

$\mathrm{Bin}(n, \lambda t/n)$, and letting $n \to \infty$,

$$P(N(t) = k) = e^{-\lambda t} \frac{(\lambda t)^k}{k!}.$$

4.6 Other discrete distributions

Geometric. The number X of trials up to and including the first success: $p(k) = (1-p)^{k-1}p$, $k \geq 1$, and $P(X > k) = (1-p)^k$. With $q = 1-p$, conditioning on the first trial gives $E[X] = p \cdot 1 + q(1 + E[X])$, so $E[X] = 1/p$; similarly $E[X^2] = p + q\,E[(1+X)^2]$ gives $E[X^2] = (2-p)/p^2$ and $\mathrm{Var}(X) = q/p^2$.

Negative binomial. The number X of trials needed for r successes: the r-th success at trial k requires $r-1$ successes in the first $k-1$ trials, so

$$p(k) = \binom{k-1}{r-1} p^r (1-p)^{k-r}, \quad k \geq r, \qquad E[X] = \frac{r}{p}, \quad \mathrm{Var}(X) = \frac{r(1-p)}{p^2}.$$

(X is a sum of r independent geometric waiting times, so Lecture 7 gives the moments at once.)

Example 4.8 (Banach match problem)

A smoker carries two boxes of N matches and each time picks a box at random. When a box is first found empty, what is the probability that the other contains exactly k matches? Call choosing the left box a "success". The left box is found empty with k matches on the right when success number $N+1$ occurs at trial $N+1+N-k$. So the probability is $2\binom{2N-k}{N}(\frac{1}{2})^{2N-k+1} = \binom{2N-k}{N}(\frac{1}{2})^{2N-k}$, the factor 2 accounting for the symmetric case.

Hypergeometric. Draw n balls without replacement from N, of which m are white; the number X of white balls drawn has

$$P(X = i) = \frac{\binom{m}{i}\binom{N-m}{n-i}}{\binom{N}{n}}, \qquad E[X] = \frac{nm}{N}, \qquad \mathrm{Var}(X) = n\frac{m}{N}\Big(1 - \frac{m}{N}\Big)\frac{N-n}{N-1}.$$

With $p = m/N$ this is the binomial mean, and the variance carries the finite-population correction $(N-n)/(N-1)$. When N is large relative to n, X is approximately $\mathrm{Bin}(n, m/N)$.

Example 4.9 (Capture–recapture)

To estimate an unknown population size N, catch and tag m animals, release them, and later catch n; suppose i are tagged. The probability of this, $P_i(N) = \binom{m}{i}\binom{N-m}{n-i}/\binom{N}{n}$, satisfies

$$\frac{P_i(N)}{P_i(N-1)} = \frac{(N-m)(N-n)}{N(N-m-n+i)} \geq 1 \iff N \leq \frac{mn}{i},$$

so the value of N that makes the observation most likely (the maximum likelihood estimate) is $\lfloor mn/i \rfloor$: the tagged fraction in the sample, i/n, is set equal to m/N.

Zeta (Zipf). $p(k) = C/k^{\alpha+1}$, $k \geq 1$, with $C = 1/\zeta(\alpha+1)$; it models heavy-tailed frequencies such as word counts.

4.7 Expected value of sums

By Proposition 4.2(c), $E[X_1+\cdots+X_n] = E[X_1]+\cdots+E[X_n]$ for *any* random variables, independent or not. Writing a count as a sum of indicators is the main trick of the course.

Example 4.10

(a) The sum of n fair dice has mean $7n/2$. (b) If trial i succeeds with probability p_i, the expected number of successes is $\sum p_i$. (c) The hypergeometric mean: with I_j the indicator that white ball j is drawn, $E[X] = \sum_{j=1}^{m} P(I_j = 1) = m \cdot n/N$ (Example 2.1(b)).

Example 4.11 (Variance of the number of successes)

With $X = \sum_i I_i$, $X^2 = \sum_i I_i + \sum_{i\neq j} I_i I_j$, so $E[X^2] = \sum_i p_i + \sum_{i\neq j} P(I_i = I_j = 1)$. For independent trials with common p, $E[X^2] = np + n(n-1)p^2$ and $\mathrm{Var}(X) = np(1-p)$. For the hypergeometric, $P(I_i = I_j = 1) = \frac{m(m-1)}{N(N-1)}$ for two specific draws $i \neq j$, which gives the variance stated above.

Lecture 5. Continuous Random Variables

5.1 Densities

X is **continuous** if there is a nonnegative f (its **probability density**) with $P(X \in B) = \int_B f(x)\,dx$ for all (reasonable) $B \subset \mathbb{R}$. Then $\int_{-\infty}^{\infty} f = 1$, $P(a \le X \le b) = \int_a^b f$, $P(X = a) = 0$, and $F(a) = \int_{-\infty}^{a} f$, so $F' = f$ wherever f is continuous. Intuitively $P(a - \frac{\epsilon}{2} \le X \le a + \frac{\epsilon}{2}) \approx \epsilon f(a)$: $f(a)$ measures how likely X is to be *near* a.

Example 5.1

(a) If $f(x) = C(4x - 2x^2)$ on $(0, 2)$ and 0 elsewhere, then $1 = C\int_0^2 (4x - 2x^2)\,dx = 8C/3$, so $C = 3/8$; by symmetry about $x = 1$, $P(X > 1) = 1/2$.
(b) If $Y = 2X$ then $F_Y(a) = F_X(a/2)$ and $f_Y(a) = \frac{1}{2} f_X(a/2)$.

5.2 Expectation and variance

$E[X] = \int x f(x)\,dx$, and $\mathrm{Var}(X) = E[(X - \mu)^2] = E[X^2] - \mu^2$ as before.

Lemma 5.1 (Tail formula)

If $Y \ge 0$, then $E[Y] = \int_0^\infty P(Y > y)\,dy$.

Proof. $\int_0^\infty P(Y > y)\,dy = \int_0^\infty \int_y^\infty f_Y(x)\,dx\,dy = \int_0^\infty \left(\int_0^x dy\right) f_Y(x)\,dx = \int_0^\infty x f_Y(x)\,dx$. ■

(The discrete analogue is $E[N] = \sum_{k \ge 1} P(N \ge k)$ for integer $N \ge 0$.)

Proposition 5.2

$E[g(X)] = \int g(x) f(x)\,dx$.

Proof. For $g \ge 0$, by the lemma, $E[g(X)] = \int_0^\infty P(g(X) > y)\,dy = \int_0^\infty \int_{x:g(x)>y} f(x)\,dx\,dy = \int \left(\int_0^{g(x)} dy\right) f(x)\,dx = \int g(x) f(x)\,dx$. In general split $g = g^+ - g^-$. ■

Example 5.2 (Breaking a stick)

A stick of length 1 is broken at a uniform point U. What is the expected length of the piece containing a fixed point $p \in [0, 1]$? The length is $L_p(U) = 1 - U$ if $U < p$ and U if $U > p$, so

$$E[L_p(U)] = \int_0^p (1 - u)\,du + \int_p^1 u\,du = \frac{1}{2} + p(1 - p).$$

It is at least $\frac{1}{2}$ and maximal for $p = \frac{1}{2}$: the piece containing p tends to be the longer one, because longer pieces are more likely to contain p (compare Example 4.3).

Example 5.3 (When to leave)

Travel time X has density f and cdf F. Arriving s minutes early costs cs; arriving s minutes

late costs ks. Leaving t minutes before the appointment costs, in expectation,

$$E[C_t(X)] = \int_0^t c(t-x)f(x)\,dx + \int_t^\infty k(x-t)f(x)\,dx, \qquad \frac{d}{dt}E[C_t(X)] = cF(t) - k\{1-F(t)\}.$$

The optimum is the quantile $F(t^*) = k/(k+c)$; for $k = c$ it is the median. The same argument shows that $E|X-a|$ is minimized at the median of X.

5.3 The uniform distribution

$X \sim U(\alpha, \beta)$ has density $1/(\beta-\alpha)$ on (α, β). Then $E[X] = (\alpha+\beta)/2$ and $\mathrm{Var}(X) = (\beta-\alpha)^2/12$.

Example 5.4 (Bertrand's paradox)

What is the probability that a "random chord" of a circle of radius r is longer than the side $\sqrt{3}r$ of the inscribed equilateral triangle? A chord at distance D from the centre has length $2\sqrt{r^2-D^2} > \sqrt{3}r$ iff $D < r/2$. (i) If $D \sim U(0,r)$ (pick a random point on a radius and draw the perpendicular chord), the answer is $1/2$. (ii) If one endpoint is fixed and the angle θ between the chord and the tangent there is $U(0,\pi)$ (equivalently, the other endpoint is uniform on the circle), the chord has length $2r\sin\theta$, which exceeds $\sqrt{3}r$ iff $\theta \in (\pi/3, 2\pi/3)$: probability $1/3$. The answers differ because "random chord" has not been defined: different experiments give different distributions.

5.4 Normal random variables

$X \sim N(\mu, \sigma^2)$ has density

$$f(x) = \frac{1}{\sqrt{2\pi}\,\sigma} e^{-(x-\mu)^2/(2\sigma^2)}, \qquad -\infty < x < \infty.$$

That it integrates to 1 follows from $I^2 = \iint e^{-(x^2+y^2)/2}\,dx\,dy = \int_0^{2\pi}\int_0^\infty e^{-r^2/2} r\,dr\,d\theta = 2\pi$ for $I = \int e^{-x^2/2}\,dx$. If $X \sim N(\mu,\sigma^2)$ then $aX+b \sim N(a\mu+b, a^2\sigma^2)$, so $Z = (X-\mu)/\sigma$ is **standard normal**, with cdf Φ and $\Phi(-x) = 1-\Phi(x)$. For Z, $E[Z] = 0$ by symmetry and $E[Z^2] = 1$ by integration by parts; hence $E[X] = \mu$, $\mathrm{Var}(X) = \sigma^2$, and

$$P(X \le a) = \Phi\Big(\frac{a-\mu}{\sigma}\Big).$$

Useful values: $P(|Z|<1) \approx 0.683$, $P(|Z|<2) \approx 0.954$, $P(|Z|<3) \approx 0.997$; $\Phi(1.645) = 0.95$, $\Phi(1.96) = 0.975$. For large x, $1-\Phi(x) \approx e^{-x^2/2}/(x\sqrt{2\pi})$.

Example 5.5 (Detecting a signal)

A bit is sent as $+2$ (for 1) or -2 (for 0); the receiver gets $R = \pm 2 + N$ with $N \sim N(0,1)$ and decodes 1 iff $R \ge 0.5$. The two error probabilities are $P(N < -1.5) = 1-\Phi(1.5) \approx 0.067$ and $P(N \ge 2.5) = 1-\Phi(2.5) \approx 0.006$.

Theorem 5.3 (De Moivre–Laplace)

If $S_n \sim \mathrm{Bin}(n,p)$, then for $a < b$, $P\Big(a \le \dfrac{S_n - np}{\sqrt{np(1-p)}} \le b\Big) \to \Phi(b) - \Phi(a)$ as $n \to \infty$.

This is a special case of the central limit theorem (Lecture 8). The approximation is good when $np(1-p) \geq 10$. Because S_n is integer-valued we use the **continuity correction**: $P(S_n = i) = P(i - \frac{1}{2} < S_n < i + \frac{1}{2})$ before standardizing. (For small p and np moderate, use the Poisson approximation instead.)

Example 5.6

(a) For 40 fair flips, $P(S = 20) \approx \Phi(0.5/\sqrt{10}) - \Phi(-0.5/\sqrt{10}) \approx 0.1256$; the exact value is $\binom{40}{20}2^{-40} \approx 0.1254$.
(b) A college admits 450 students, each enrolling with probability 0.3 independently, and has room for 150. $S \sim \text{Bin}(450, 0.3)$ has mean 135 and variance 94.5, so $P(S \geq 151) = P(S \geq 150.5) \approx 1 - \Phi(15.5/\sqrt{94.5}) = 1 - \Phi(1.59) \approx 0.055$ (exact: 0.057).

5.5 Exponential random variables

$X \sim \text{Exp}(\lambda)$ has $f(x) = \lambda e^{-\lambda x}$ and $P(X > x) = e^{-\lambda x}$ for $x \geq 0$; $E[X] = 1/\lambda$, $\text{Var}(X) = 1/\lambda^2$. In a Poisson process with rate λ (Example 4.7), the time to the first event exceeds t iff $N(t) = 0$, which has probability $e^{-\lambda t}$, so waiting times are exponential.

Proposition 5.4 (Memorylessness)

$P(X > s + t \mid X > t) = P(X > s)$ for all $s, t \geq 0$ holds for $X \sim \text{Exp}(\lambda)$, and the exponential distributions are the only continuous distributions on $[0, \infty)$ with this property.

Proof. For the exponential, $e^{-\lambda(s+t)}/e^{-\lambda t} = e^{-\lambda s}$. Conversely, $\bar{F}(x) = P(X > x)$ satisfies $\bar{F}(s+t) = \bar{F}(s)\bar{F}(t)$. Then $\bar{F}(m/n) = \bar{F}(1/n)^m$ and $\bar{F}(1) = \bar{F}(1/n)^n$, so $\bar{F}(x) = \bar{F}(1)^x$ for rational x, and by right continuity for all $x \geq 0$; i.e. $\bar{F}(x) = e^{-\lambda x}$ with $\lambda = -\log \bar{F}(1)$. ∎

Example 5.7 (Two clerks)

A customer enters a post office where clerks are serving Ms Jones and Mr Brown, and will be served as soon as either finishes. Service times are independent exponentials with the same rate. What is the probability that the customer is the last of the three to leave? When the first of Jones and Brown leaves, the other's remaining service time is again exponential with the same rate, by memorylessness, so the newcomer and the remaining customer are in a symmetric position: the probability is 1/2.

Hazard rate. For a positive continuous lifetime X with cdf F, the **hazard (failure) rate** is $\lambda(t) = f(t)/\bar{F}(t)$; $\lambda(t)\,dt$ is the probability that an item of age t fails in the next dt. Since $\lambda(t) = -\frac{d}{dt}\log \bar{F}(t)$,

$$\bar{F}(t) = \exp\Big\{-\int_0^t \lambda(s)\,ds\Big\},$$

so λ determines the distribution. A constant hazard characterizes the exponential; increasing hazards describe ageing.

Example 5.8

If at every age a smoker's death rate is twice a non-smoker's, the probability that a smoker of age A survives to age B is $\exp\{-2\int_A^B \lambda_N\}$, the *square* of a non-smoker's survival probability. If the latter is 0.76 from age 50 to 60, the smoker's is 0.578.

5.6 Other continuous distributions

Gamma. $X \sim \text{Gamma}(\alpha, \lambda)$ has $f(x) = \lambda e^{-\lambda x}(\lambda x)^{\alpha-1}/\Gamma(\alpha)$, $x > 0$, where $\Gamma(\alpha) = \int_0^\infty e^{-y} y^{\alpha-1}\, dy$ satisfies $\Gamma(\alpha) = (\alpha-1)\Gamma(\alpha-1)$ (integration by parts), $\Gamma(n) = (n-1)!$, and $\Gamma(\frac{1}{2}) = \sqrt{\pi}$. Then $E[X] = \alpha/\lambda$ and $\text{Var}(X) = \alpha/\lambda^2$. The time T_n of the n-th event of a Poisson process is $\text{Gamma}(n, \lambda)$, because

$$P(T_n \le t) = P(N(t) \ge n) = \sum_{j \ge n} e^{-\lambda t} \frac{(\lambda t)^j}{j!},$$

and differentiating in t the sum telescopes to $\lambda e^{-\lambda t}(\lambda t)^{n-1}/(n-1)!$. The **chi-squared** distribution with n degrees of freedom is $\text{Gamma}(n/2, 1/2)$.

Weibull. $F(x) = 1 - \exp\{-((x-\nu)/\alpha)^\beta\}$ for $x > \nu$; it models lifetimes (weakest-link failures) and has hazard proportional to $(x-\nu)^{\beta-1}$.

Cauchy. $f(x) = \frac{1}{\pi}\frac{1}{1+(x-\theta)^2}$. A flashlight at distance 1 from a wall is spun so that its angle θ to the perpendicular is $U(-\pi/2, \pi/2)$; the point lit on the wall is $X = \tan\theta$, and $P(X \le x) = P(\theta \le \arctan x) = \frac{1}{2} + \frac{1}{\pi}\arctan x$, the standard Cauchy. Its mean does not exist, since $\int |x|/(1+x^2)\, dx = \infty$.

Beta. $f(x) = x^{a-1}(1-x)^{b-1}/B(a,b)$ on $(0,1)$, with $B(a,b) = \Gamma(a)\Gamma(b)/\Gamma(a+b)$, $E[X] = a/(a+b)$ and $\text{Var}(X) = ab/\{(a+b)^2(a+b+1)\}$. It models a random proportion; $a = b = 1$ is uniform.

Pareto. If $X \sim \text{Exp}(\lambda)$ and $a > 0$, then $Y = ae^X$ has $P(Y > y) = (a/y)^\lambda$ for $y \ge a$. Given $Y > y_0$, Y is again Pareto with a replaced by y_0: incomes above a threshold look like rescaled incomes.

5.7 The distribution of a function of a random variable

The general method is to compute the cdf of $Y = g(X)$ and differentiate.

Theorem 5.5

If g is strictly monotone and differentiable on the range of X, then $Y = g(X)$ has density

$$f_Y(y) = f_X(g^{-1}(y)) \left| \frac{d}{dy} g^{-1}(y) \right|$$

for y in the range of g, and 0 otherwise.

Proof. For increasing g, $F_Y(y) = P(X \le g^{-1}(y)) = F_X(g^{-1}(y))$; differentiate. For decreasing g, $F_Y(y) = 1 - F_X(g^{-1}(y))$, and the derivative of g^{-1} is negative. ■

Example 5.9

(a) $Y = X^2$ (not monotone): for $y > 0$, $F_Y(y) = F_X(\sqrt{y}) - F_X(-\sqrt{y})$, so $f_Y(y) = \frac{1}{2\sqrt{y}}\{f_X(\sqrt{y}) + f_X(-\sqrt{y})\}$. With $X = Z$ standard normal this is $\frac{1}{\sqrt{2\pi y}} e^{-y/2}$, the chi-squared density with one degree of freedom (and shows $\Gamma(\frac{1}{2}) = \sqrt{\pi}$).
(b) $|X|$ has density $f_X(y) + f_X(-y)$ for $y > 0$.
(c) Lognormal: if $X \sim N(\mu, \sigma^2)$, $Y = e^X$ has $f_Y(y) = \frac{1}{\sqrt{2\pi}\sigma y} \exp\{-(\log y - \mu)^2/(2\sigma^2)\}$, $y > 0$. It models prices, where daily ratios multiply.

(d) If F is continuous and strictly increasing, then $F(X) \sim U(0,1)$ and $F^{-1}(U)$ has cdf F for $U \sim U(0,1)$; this is how random variables are simulated.

Lecture 6. Jointly Distributed Random Variables

6.1 Joint distribution functions

The **joint cdf** of X, Y is $F(a,b) = P(X \le a, Y \le b)$; the marginals are $F_X(a) = F(a,\infty)$ and $F_Y(b) = F(\infty, b)$, and

$$P(a_1 < X \le a_2,\ b_1 < Y \le b_2) = F(a_2,b_2) - F(a_1,b_2) - F(a_2,b_1) + F(a_1,b_1).$$

For discrete variables the joint pmf is $p(x,y) = P(X=x, Y=y)$, and $p_X(x) = \sum_y p(x,y)$. X and Y are **jointly continuous** with joint density f if $P((X,Y) \in C) = \iint_C f(x,y)\,dx\,dy$; then $f = \partial^2 F/\partial a\,\partial b$, and $f_X(x) = \int f(x,y)\,dy$. All of this extends to n variables.

Example 6.1

Let $f(x,y) = 2e^{-x}e^{-2y}$ for $x, y > 0$. Then $P(X > 1, Y < 1) = \int_0^1 \int_1^\infty 2e^{-x}e^{-2y}\,dx\,dy = e^{-1}(1-e^{-2}) \approx 0.318$, and

$$P(X < Y) = \int_0^\infty \int_0^y 2e^{-x}e^{-2y}\,dx\,dy = \int_0^\infty 2e^{-2y}(1-e^{-y})\,dy = 1 - \tfrac{2}{3} = \tfrac{1}{3}.$$

Example 6.2 (A random point in a disc)

(X,Y) is uniform on the disc of radius R: $f = 1/(\pi R^2)$ there. The marginal is $f_X(x) = \frac{2\sqrt{R^2-x^2}}{\pi R^2}$ for $|x| \le R$, which is not uniform. The distance D to the centre has $P(D \le a) = \pi a^2/(\pi R^2) = a^2/R^2$, so $f_D(a) = 2a/R^2$ and $E[D] = 2R/3$.

Example 6.3 (Multinomial distribution)

In n independent trials, each resulting in outcome i with probability p_i ($\sum_1^r p_i = 1$), the counts $X_1, \dots, X_r$ satisfy

$$P(X_1 = n_1, \dots, X_r = n_r) = \frac{n!}{n_1! \cdots n_r!} p_1^{n_1} \cdots p_r^{n_r}, \qquad \sum n_i = n,$$

since each sequence with these counts has probability $\prod p_i^{n_i}$ and there are $\binom{n}{n_1,\dots,n_r}$ such sequences (Lecture 1). Each X_i is marginally $\mathrm{Bin}(n, p_i)$.

6.2 Independent random variables

X and Y are **independent** if $P(X \in A, Y \in B) = P(X \in A)P(Y \in B)$ for all A, B; equivalently $F(a,b) = F_X(a)F_Y(b)$, or $p(x,y) = p_X(x)p_Y(y)$, or $f(x,y) = f_X(x)f_Y(y)$.

Proposition 6.1 (Factorization)

Jointly continuous (or discrete) X, Y are independent iff the joint density (pmf) factors as $f(x,y) = h(x)g(y)$ for all x, y.

Proof. If $f = hg$, put $C_1 = \int h$, $C_2 = \int g$; then $C_1C_2 = \iint f = 1$, $f_X = hC_2$ and $f_Y = gC_1$, so $f_X f_Y = hg\,C_1C_2 = f$. ∎

The factorization must hold on the whole plane: $f = 24xy$ on $\{x, y > 0, x + y < 1\}$ looks like a product but is not, because the support is not a rectangle.

Example 6.4 (Thinning a Poisson count)

The number N of customers on a day is Poisson(λ), and each is independently a man with probability p. Let X and Y be the numbers of men and women. Then

$$P(X=i,Y=j)=e^{-\lambda}\frac{\lambda^{i+j}}{(i+j)!}\binom{i+j}{i}p^iq^j=e^{-\lambda p}\frac{(\lambda p)^i}{i!}\cdot e^{-\lambda q}\frac{(\lambda q)^j}{j!}:$$

X and Y are *independent* Poisson variables with means λp and λq. (With a fixed number n of customers, they would be strongly dependent.)

Example 6.5 (Meeting)

Two people arrive independently and uniformly between 12:15 and 12:45 and each waits 10 minutes. Measuring in minutes past 12:15, they meet iff $|X-Y|\le 10$, so $P=1-(20/30)^2=5/9$ (the two corner triangles of the square have total area 20^2).

Example 6.6 (Buffon's needle)

A needle of length L is dropped on a floor ruled with parallel lines $D\ge L$ apart. Let $X\sim U(0,D/2)$ be the distance from the needle's centre to the nearest line and $\theta\sim U(0,\pi/2)$ the acute angle with the lines, independent. The needle crosses a line iff $X<\frac{L}{2}\cos\theta$, so

$$P(\text{cross})=\frac{4}{\pi D}\int_0^{\pi/2}\int_0^{(L/2)\cos\theta}dx\,d\theta=\frac{2L}{\pi D}.$$

Dropping many needles thus estimates π.

Example 6.7 (Characterizing the normal)

Darts land at (X,Y) with X,Y independent with a common differentiable density f, and the joint density depends only on the distance from the target: $f(x)f(y)=g(x^2+y^2)$. Differentiating in x gives $f'(x)f(y)=2xg'(x^2+y^2)$, so $\frac{f'(x)}{2xf(x)}=\frac{g'(x^2+y^2)}{g(x^2+y^2)}$ for all x,y. The left side does not depend on y, and the right side then shows it does not depend on x; hence $f'(x)/f(x)=cx$, $f(x)=ke^{cx^2/2}$ with $c<0$: X is normal with mean 0.

6.3 Sums of independent random variables

If X,Y are independent with densities f_X,f_Y, then

$$F_{X+Y}(a)=\int F_X(a-y)f_Y(y)\,dy,\qquad f_{X+Y}(a)=\int f_X(a-y)f_Y(y)\,dy,$$

the **convolution** of f_X and f_Y.

Example 6.8 (Sums of uniforms)

For $X,Y\sim U(0,1)$ independent, $f_{X+Y}(a)=a$ on $[0,1]$ and $2-a$ on $[1,2]$ (triangular). More generally, if S_n is a sum of n independent uniforms, induction on n gives $P(S_n\le x)=x^n/n!$ for $0\le x\le 1$. Let $N=\min\{n:S_n>1\}$. Then $P(N>n)=P(S_n\le 1)=1/n!$, so

$$E[N]=\sum_{n\ge 0}P(N>n)=\sum_{n\ge 0}\frac{1}{n!}=e.$$

Proposition 6.2

For independent variables: (a) $\text{Gamma}(s,\lambda)+\text{Gamma}(t,\lambda) \sim \text{Gamma}(s+t,\lambda)$; in particular a sum of n independent $\text{Exp}(\lambda)$ is $\text{Gamma}(n,\lambda)$ and $Z_1^2+\cdots+Z_n^2 \sim \chi^2_n$. (b) $\sum N(\mu_i,\sigma_i^2) \sim N(\sum \mu_i, \sum \sigma_i^2)$. (c) $\text{Poisson}(\lambda_1)+\text{Poisson}(\lambda_2) \sim \text{Poisson}(\lambda_1+\lambda_2)$. (d) $\text{Bin}(n,p)+\text{Bin}(m,p) \sim \text{Bin}(n+m,p)$.

Proof. (a) The convolution is $\frac{\lambda e^{-\lambda a}}{\Gamma(s)\Gamma(t)}\lambda^{s+t-1}\int_0^a y^{t-1}(a-y)^{s-1}\,dy$; substituting $y=ax$ shows it is $Ca^{s+t-1}e^{-\lambda a}$, a gamma density, and C is then forced. (b) It suffices to treat $N(0,\sigma^2)+N(0,1)$, a direct completion of the square; Lecture 7 gives a one-line proof with moment generating functions. (c) $P(X+Y=n)=\sum_k e^{-\lambda_1}\frac{\lambda_1^k}{k!}e^{-\lambda_2}\frac{\lambda_2^{n-k}}{(n-k)!} = e^{-(\lambda_1+\lambda_2)}\frac{(\lambda_1+\lambda_2)^n}{n!}$ by the binomial theorem. (d) Count successes in $n+m$ trials, or use Vandermonde. ∎

Example 6.9

A team plays 26 games against class A teams, winning each with probability 0.4, and 18 against class B, winning each with probability 0.7, all independently. Its number of wins has mean $26(0.4)+18(0.7)=23$ and variance $26(0.24)+18(0.21)=10.02$; it is not binomial, but the normal approximation applies: $P(\text{wins} \ge 25) \approx 1-\Phi(1.5/\sqrt{10.02}) \approx 0.32$.

6.4 Conditional distributions

For discrete variables, $p_{X|Y}(x \mid y) = p(x,y)/p_Y(y)$ when $p_Y(y)>0$. For jointly continuous ones, the conditional density is

$$f_{X|Y}(x \mid y) = \frac{f(x,y)}{f_Y(y)}, \qquad P(X \in A \mid Y=y) = \int_A f_{X|Y}(x \mid y)\,dx,$$

which is motivated by $P(x<X<x+dx \mid y<Y<y+dy) \approx f(x,y)\,dx\,dy/(f_Y(y)\,dy)$. Independence means the conditional law does not depend on y.

Example 6.10

(a) If X,Y are independent Poisson with means λ_1,λ_2, then given $X+Y=n$, $X \sim \text{Bin}(n,\lambda_1/(\lambda_1+\lambda_2))$. Likewise, given $X+Y=n$ for independent $\text{Bin}(n_1,p)$ and $\text{Bin}(n_2,p)$, X is hypergeometric, free of p.
(b) In a multinomial, given $X_r=n_r$, the remaining counts are multinomial with $n-n_r$ trials and probabilities $p_i/(1-p_r)$.
(c) If $f(x,y)=e^{-x/y}e^{-y}/y$ for $x,y>0$, then $f_{X|Y}(x \mid y)=\frac{1}{y}e^{-x/y}$: given $Y=y$, X is exponential with mean y, and $P(X>1 \mid Y=y)=e^{-1/y}$.

Example 6.11 (The bivariate normal)

(X,Y) is bivariate normal with parameters $\mu_x,\mu_y,\sigma_x,\sigma_y,\rho$ ($|\rho|<1$) if

$$f(x,y)=\frac{1}{2\pi\sigma_x\sigma_y\sqrt{1-\rho^2}}\exp\Big\{-\frac{1}{2(1-\rho^2)}\big[u^2-2\rho uv+v^2\big]\Big\}, \quad u=\frac{x-\mu_x}{\sigma_x},\ v=\frac{y-\mu_y}{\sigma_y}.$$

Completing the square in u shows that, given $Y=y$, X is normal with

$$E[X \mid Y=y]=\mu_x+\rho\frac{\sigma_x}{\sigma_y}(y-\mu_y), \qquad \text{Var}(X \mid Y=y)=\sigma_x^2(1-\rho^2),$$

and integrating out x shows $Y \sim N(\mu_y,\sigma_y^2)$. X and Y are independent iff $\rho=0$ (Lecture 7 shows ρ is the correlation).

Example 6.12 (Mixed variables: beta–binomial)

Let $P \sim U(0,1)$ and, given $P = p$, let $N \sim \mathrm{Bin}(n,p)$. Then $P(N=k) = \int_0^1 \binom{n}{k} p^k (1-p)^{n-k}\,dp = \binom{n}{k} B(k+1, n-k+1) = \frac{1}{n+1}$: the number of successes is uniform on $\{0,\dots,n\}$. Given $N=k$, the density of P is proportional to $p^k(1-p)^{n-k}$: $P \mid N=k \sim \mathrm{Beta}(k+1, n-k+1)$. Compare Laplace's rule of succession (Example 3.12).

6.5 Order statistics

Let $X_1,\dots,X_n$ be i.i.d. with density f and cdf F, and $X_{(1)} \le \cdots \le X_{(n)}$ their ordered values. Each of the $n!$ assignments of the X_i to the ordered values is equally likely, so

$$f_{X_{(1)},\dots,X_{(n)}}(x_1,\dots,x_n) = n!\, f(x_1)\cdots f(x_n), \qquad x_1 < \cdots < x_n.$$

For one order statistic: $X_{(j)}$ is near x when $j-1$ observations are below x, one is near x, and $n-j$ are above, so

$$f_{X_{(j)}}(x) = \frac{n!}{(j-1)!\,(n-j)!} F(x)^{j-1}\big(1-F(x)\big)^{n-j} f(x).$$

Also $P(X_{(j)} \le x) = \sum_{k=j}^{n} \binom{n}{k} F(x)^k (1-F(x))^{n-k}$ (at least j observations below x). For uniforms, $X_{(j)} \sim \mathrm{Beta}(j, n-j+1)$, so $E[X_{(j)}] = j/(n+1)$.

Example 6.13

(a) Three people are placed independently and uniformly on a road of length 1. The probability that no two are within distance $d \le \frac{1}{2}$ is

$$3! \int_{x_1 < x_2 - d,\ x_2 < x_3 - d} dx_1\,dx_2\,dx_3 = (1-2d)^3,$$

by substituting $y_1 = x_1$, $y_2 = x_2 - d$, $y_3 = x_3 - 2d$, which maps the region onto the ordered simplex in $[0, 1-2d]^3$.

(b) The sample median of $2n+1$ uniforms has density $\frac{(2n+1)!}{n!\,n!} x^n (1-x)^n$, and the range $R = X_{(n)} - X_{(1)}$ has $P(R \le a) = n \int f(x)\{F(x+a) - F(x)\}^{n-1}\,dx$; for uniforms, $f_R(a) = n(n-1)a^{n-2}(1-a)$.

6.6 Functions of random variables: the Jacobian method

Theorem 6.3

Let (X_1, X_2) have density f, and let $y_1 = g_1(x_1,x_2)$, $y_2 = g_2(x_1,x_2)$ be a one-to-one differentiable map with inverse $x_i = h_i(y_1,y_2)$ and nonzero Jacobian $J(x_1,x_2) = \det(\partial g_i/\partial x_j)$. Then $Y_i = g_i(X_1,X_2)$ have joint density

$$f_{Y_1,Y_2}(y_1,y_2) = f\big(h_1(y),h_2(y)\big)\,\big|J(h_1(y),h_2(y))\big|^{-1}.$$

This is the change of variables formula for double integrals; it holds in n dimensions as well.

Example 6.14 (Polar coordinates of normals)

Let X, Y be independent $N(0,1)$, and $R = \sqrt{X^2+Y^2}$, $\Theta = \arctan(Y/X) \in [0, 2\pi)$. The

Jacobian of $(x,y)\mapsto(r,\theta)$ is $1/r$, so

$$f_{R,\Theta}(r,\theta)=\frac{1}{2\pi}e^{-r^2/2}\,r=\frac{1}{2\pi}\cdot re^{-r^2/2}:$$

R and Θ are independent, Θ is uniform, and $R^2\sim\text{Exp}(1/2)=\chi^2_2$. Reversing this gives the Box–Muller method: $X=\sqrt{-2\log U_1}\cos(2\pi U_2)$ and $Y=\sqrt{-2\log U_1}\sin(2\pi U_2)$ are independent standard normals.

Example 6.15 (Gamma and beta)

Let $X\sim\text{Gamma}(\alpha,\lambda)$ and $Y\sim\text{Gamma}(\beta,\lambda)$ be independent, $U=X+Y$, $V=X/(X+Y)$. Then $x=uv$, $y=u(1-v)$, $|J|^{-1}=u$, and

$$f_{U,V}(u,v)=\frac{\lambda e^{-\lambda u}(\lambda u)^{\alpha+\beta-1}}{\Gamma(\alpha+\beta)}\cdot\frac{\Gamma(\alpha+\beta)}{\Gamma(\alpha)\Gamma(\beta)}v^{\alpha-1}(1-v)^{\beta-1}.$$

So $U\sim\text{Gamma}(\alpha+\beta,\lambda)$, $V\sim\text{Beta}(\alpha,\beta)$, and they are independent; as a by-product $B(\alpha,\beta)=\Gamma(\alpha)\Gamma(\beta)/\Gamma(\alpha+\beta)$. For instance, if two jobs with exponential durations are done in sequence, the fraction of time spent on the first is uniform and independent of the total.

Example 6.16 (Spacings of exponentials)

For i.i.d. $\text{Exp}(\lambda)$ variables $X_1,\dots,X_n$, the spacings $Y_1=X_{(1)}$, $Y_j=X_{(j)}-X_{(j-1)}$ are independent with $Y_j\sim\text{Exp}((n-j+1)\lambda)$. Indeed the map from ordered values to spacings has Jacobian 1, and $\sum_j x_{(j)}=\sum_j(n-j+1)y_j$, so the joint density $n!\lambda^n e^{-\lambda\sum x_{(j)}}$ factors. Intuitively, after the first failure the $n-1$ survivors start afresh by memorylessness. Hence $E[X_{(n)}]=\frac{1}{\lambda}\sum_{k=1}^{n}\frac{1}{k}$.

6.7 Exchangeable random variables

$X_1,\dots,X_n$ are **exchangeable** if $(X_{i_1},\dots,X_{i_n})$ has the same joint law for every permutation. I.i.d. variables are exchangeable, but so are dependent ones such as the indicators X_i that the i-th ball drawn without replacement is white: any particular sequence with k whites has probability $\frac{m(m-1)\cdots(m-k+1)\,(N-m)\cdots(N-m-n+k+1)}{N(N-1)\cdots(N-n+1)}$, which depends only on k. So every draw is white with probability m/N, as in Example 2.1(b). Pólya's urn (draw a ball, return it with another of the same colour) also gives exchangeable draws.

Lecture 7. Properties of Expectation

7.1 Expectation of sums

For jointly distributed X, Y, $E[g(X,Y)] = \sum_x \sum_y g(x,y)p(x,y)$ or $\iint g(x,y)f(x,y)\,dx\,dy$ (proved as Proposition 5.2). Taking $g = x + y$ gives linearity, $E[X+Y] = E[X] + E[Y]$, and by induction $E[\sum X_i] = \sum E[X_i]$, with no independence needed. Also $X \geq Y$ implies $E[X] \geq E[Y]$.

Example 7.1

(a) An accident at $X \sim U(0,L)$ and an ambulance at independent $Y \sim U(0,L)$: $E|X-Y| = \frac{1}{L^2}\int_0^L \int_0^L |x-y|\,dx\,dy = \frac{L}{3}$.
(b) Boole's inequality: if X counts how many of $A_1, \ldots, A_n$ occur and $Y = \mathbb{1}\{X \geq 1\}$, then $X \geq Y$, so $P(\bigcup A_i) = E[Y] \leq E[X] = \sum P(A_i)$.
(c) The sample mean $\bar{X} = \frac{1}{n}\sum X_i$ of identically distributed variables has $E[\bar{X}] = \mu$.

Example 7.2 (Indicator calculations)

(a) Matches: $E[\#\text{matches}] = N \cdot \frac{1}{N} = 1$ for every N.
(b) Coupon collecting with N types: once i types are held, the number of further purchases until a new type is geometric with success probability $(N-i)/N$. So the total T has

$$E[T] = \sum_{i=0}^{N-1} \frac{N}{N-i} = N\Big(1 + \frac{1}{2} + \cdots + \frac{1}{N}\Big) \approx N \log N.$$

The expected number of distinct types in n purchases is $N\{1-(1-1/N)^n\}$, by summing indicators of "type j appears".
(c) Runs: in a random arrangement of n ones and m zeros, the expected number of runs of ones is $\frac{n(m+1)}{n+m}$. Let I_i be the indicator that a run of ones starts at position i; $P(I_1 = 1) = \frac{n}{n+m}$, and for $i > 1$, $P(I_i = 1) = P(\text{position } i-1 \text{ is } 0,\ i \text{ is } 1) = \frac{mn}{(n+m)(n+m-1)}$.
(d) Quicksort sorts n distinct values by choosing a random pivot, comparing all others with it, and recursing on both sides. The i-th and j-th smallest values ($i < j$) are compared iff one of them is the first pivot chosen among the $j-i+1$ values between them inclusive, which has probability $2/(j-i+1)$. Hence

$$E[\#\text{comparisons}] = \sum_{i<j} \frac{2}{j-i+1} = 2(n+1)\sum_{k=1}^{n} \frac{1}{k} - 4n \approx 2n \log n.$$

Proposition 7.1 (Moments of the number of events)

If X is the number of the events $A_1, \ldots, A_n$ that occur, then

$$E\Big[\binom{X}{k}\Big] = \sum_{i_1 < \cdots < i_k} P(A_{i_1} \cdots A_{i_k}).$$

Proof. $\binom{X}{k}$ is the number of k-subsets of the events that all occur, i.e. $\sum_{i_1<\cdots<i_k} I_{i_1} \cdots I_{i_k}$; take expectations. ∎

With $k = 2$, $E[X^2] = E[X] + 2\sum_{i<j} P(A_iA_j)$. For matches, $P(A_iA_j) = \frac{1}{N(N-1)}$, so $E[X(X-1)] = 1$ and $\mathrm{Var}(X) = 1$, as for Poisson(1).

7.2 Covariance and variance of sums

Proposition 7.2

If X, Y are independent, $E[g(X)h(Y)] = E[g(X)]E[h(Y)]$.

Proof. $\iint g(x)h(y)f_X(x)f_Y(y)\,dx\,dy$ factors. ∎

$\mathrm{Cov}(X,Y) = E[(X-\mu_X)(Y-\mu_Y)] = E[XY] - E[X]E[Y]$. Independence implies Cov $= 0$, but not conversely: for X uniform on $\{-1,0,1\}$ and $Y = X^2$, $\mathrm{Cov}(X,Y) = E[X^3] = 0$.

Proposition 7.3

Covariance is symmetric and bilinear: $\mathrm{Cov}(\sum_i a_iX_i, \sum_j b_jY_j) = \sum_{i,j} a_ib_j \mathrm{Cov}(X_i,Y_j)$, and $\mathrm{Cov}(X,X) = \mathrm{Var}(X)$. Hence

$$\mathrm{Var}\Big(\sum_i X_i\Big) = \sum_i \mathrm{Var}(X_i) + 2\sum_{i<j} \mathrm{Cov}(X_i,X_j),$$

which is $\sum \mathrm{Var}(X_i)$ for pairwise independent variables.

Example 7.3

(a) For i.i.d. X_i with variance σ^2, $\mathrm{Var}(\bar X) = \sigma^2/n$, and since $\sum(X_i-\bar X)^2 = \sum(X_i-\mu)^2 - n(\bar X-\mu)^2$,

$$E[S^2] = \frac{1}{n-1}(n\sigma^2 - n\cdot\sigma^2/n) = \sigma^2, \qquad S^2 = \frac{1}{n-1}\sum(X_i-\bar X)^2.$$

(b) Hypergeometric variance: with X_i the indicator that the i-th draw is white, $\mathrm{Cov}(X_i,X_j) = \frac{m(m-1)}{N(N-1)} - \frac{m^2}{N^2} = -\frac{m(N-m)}{N^2(N-1)}$ for $i \neq j$, which with n draws gives $\mathrm{Var}(X) = n\frac{m}{N}(1-\frac{m}{N})\frac{N-n}{N-1}$.
(c) Multinomial: if N_i is the number of trials with outcome i, then $\mathrm{Cov}(N_i,N_j) = -np_ip_j$ for $i \neq j$: write $N_i = \sum_k I_i(k)$; $I_i(k)I_j(k) = 0$ and trials are independent, so $\mathrm{Cov}(N_i,N_j) = \sum_k \mathrm{Cov}(I_i(k),I_j(k)) = n(0-p_ip_j)$.

The **correlation** is $\rho(X,Y) = \mathrm{Cov}(X,Y)/\sqrt{\mathrm{Var}(X)\,\mathrm{Var}(Y)}$.

Proposition 7.4

$-1 \le \rho(X,Y) \le 1$, with $\rho = \pm 1$ iff $Y = a + bX$ with probability 1 for some b of the same sign.

Proof. With σ_X^2, σ_Y^2 the variances, $0 \le \mathrm{Var}(X/\sigma_X \pm Y/\sigma_Y) = 2(1\pm\rho)$. Equality means $X/\sigma_X \pm Y/\sigma_Y$ is constant. ∎

ρ measures *linear* dependence only. For the bivariate normal of Example 6.11, $\rho(X,Y) = \rho$.

7.3 Conditional expectation

$E[X \mid Y=y] = \sum_x x p_{X|Y}(x \mid y)$ or $\int x f_{X|Y}(x \mid y)\,dx$. $E[X \mid Y]$ denotes the random variable $g(Y)$ where $g(y) = E[X \mid Y=y]$.

Theorem 7.5 (Tower property)

$E[X] = E\big[E[X \mid Y]\big]$, i.e. $E[X] = \sum_y E[X \mid Y = y]P(Y = y)$ or $\int E[X \mid Y = y]f_Y(y)\,dy$.

Proof. Discrete case: $\sum_y \sum_x x p_{X|Y}(x \mid y)p_Y(y) = \sum_x \sum_y x p(x, y) = E[X]$. ■

Example 7.4 (The trapped miner)

A miner chooses at random among three doors: door 1 leads to safety in 3 hours, doors 2 and 3 lead back to the mine after 5 and 7 hours. With X the time to safety and Y the first door, $E[X] = \frac{1}{3}\{3 + (5 + E[X]) + (7 + E[X])\}$, so $E[X] = 15$.

Example 7.5 (Random sums)

If N is a nonnegative integer variable independent of i.i.d. $X_1, X_2, \ldots$, then $E[\sum_{i=1}^N X_i \mid N = n] = nE[X]$, so

$$E\Big[\sum_{i=1}^{N} X_i\Big] = E[N]E[X].$$

(The expected spending in a store is the expected number of customers times the mean spending.)

Example 7.6 (Waiting for k consecutive successes)

In independent trials with success probability p, let N_k be the number of trials until k consecutive successes. Conditioning on N_{k-1}: one more trial either succeeds (done) or fails (start over), so $E[N_k \mid N_{k-1}] = N_{k-1} + 1 + (1 - p)E[N_k]$, giving $E[N_k] = \frac{1}{p} + \frac{E[N_{k-1}]}{p}$ and

$$E[N_k] = \frac{1}{p} + \frac{1}{p^2} + \cdots + \frac{1}{p^k}.$$

For a fair coin, $E[N_2] = 6$ and $E[N_3] = 14$.

Probabilities by conditioning. With $X = I_E$, the tower property reads $P(E) = E[P(E \mid Y)]$. For instance, for independent continuous X, Y, $P(X < Y) = \int F_X(y)f_Y(y)\,dy$, and $P(X + Y < a) = \int F_X(a - y)f_Y(y)\,dy$ is the convolution formula again.

Example 7.7 (The best-prize problem)

n prizes of distinct values appear in random order; each must be accepted or rejected on the spot, knowing only the relative ranks so far. Strategy: reject the first k, then accept the first prize better than all earlier ones. Let X be the position of the best prize. If $X = i > k$, we win iff the best of the first $i - 1$ prizes is among the first k, which has probability $k/(i - 1)$. So

$$P_k(\text{best}) = \frac{1}{n}\sum_{i=k+1}^{n} \frac{k}{i-1} \approx \frac{k}{n}\log\frac{n}{k}.$$

Maximizing $x\log(1/x)$ over $x = k/n$ gives $k \approx n/e$, and the success probability tends to $1/e \approx 0.368$ however large n is.

Example 7.8 (Uniform sums revisited)

Let $m(x) = E[N(x)]$, where $N(x) = \min\{n : U_1 + \cdots + U_n > x\}$ for $0 \le x \le 1$. Conditioning on $U_1 = y$: if $y > x$ then $N = 1$, otherwise $N = 1 + N'(x-y)$. So $m(x) = 1 + \int_0^x m(x-y)\,dy = 1 + \int_0^x m(u)\,du$, whence $m' = m$, $m(0) = 1$, and $m(x) = e^x$. At $x = 1$ this recovers Example 6.8.

Proposition 7.6 (Conditional variance formula)

With $\operatorname{Var}(X \mid Y) = E[(X - E[X \mid Y])^2 \mid Y]$,

$$\operatorname{Var}(X) = E\big[\operatorname{Var}(X \mid Y)\big] + \operatorname{Var}\big(E[X \mid Y]\big).$$

Proof. $E[\operatorname{Var}(X \mid Y)] = E[X^2] - E[(E[X \mid Y])^2]$ and $\operatorname{Var}(E[X \mid Y]) = E[(E[X \mid Y])^2] - (E[X])^2$; add. ∎

Example 7.9

(a) For the random sum $S = \sum_1^N X_i$: $E[S \mid N] = N\mu$ and $\operatorname{Var}(S \mid N) = N\sigma^2$, so $\operatorname{Var}(S) = E[N]\sigma^2 + \mu^2 \operatorname{Var}(N)$. With $N \sim \text{Poisson}(\lambda)$, $\operatorname{Var}(S) = \lambda E[X^2]$.
(b) Geometric variance: if Y indicates success on the first trial, $E[N \mid Y]$ is 1 or $1 + 1/p$ and $\operatorname{Var}(N \mid Y)$ is 0 or $\operatorname{Var}(N)$; solving gives $\operatorname{Var}(N) = (1-p)/p^2$.

7.4 Conditional expectation and prediction

Proposition 7.7

$E[(Y - g(X))^2] \ge E[(Y - E[Y \mid X])^2]$ for every function g: $E[Y \mid X]$ is the best predictor of Y from X in mean square error.

Proof. Given X, $E[(Y-c)^2 \mid X]$ is minimized by $c = E[Y \mid X]$ (variance $\le$ any second moment about c); take expectations. ∎

If we restrict to linear predictors $a + bX$, minimizing $E[(Y - a - bX)^2]$ gives

$$b = \rho\frac{\sigma_Y}{\sigma_X}, \qquad a = \mu_Y - b\mu_X, \qquad \text{MSE} = \sigma_Y^2(1-\rho^2).$$

For the bivariate normal $E[Y \mid X]$ is linear, so the best linear predictor is the best predictor. Galton's observation that sons of tall fathers are on average less tall ("regression to the mean") is the factor $|\rho| < 1$.

7.5 Moment generating functions

The **moment generating function** of X is $M(t) = E[e^{tX}]$. When finite near 0, $M^{(n)}(0) = E[X^n]$ (differentiate under the expectation).

Distribution	$M(t)$	valid for
Bin(n,p)	$(pe^t+1-p)^n$	all t
Poisson(λ)	$\exp\{\lambda(e^t-1)\}$	all t
Geometric(p)	$pe^t/\{1-(1-p)e^t\}$	$(1-p)e^t<1$
Exp(λ)	$\lambda/(\lambda-t)$	$t<\lambda$
Gamma(α,λ)	$\{\lambda/(\lambda-t)\}^\alpha$	$t<\lambda$
$N(\mu,\sigma^2)$	$\exp\{\mu t+\sigma^2t^2/2\}$	all t

For the normal: $E[e^{tZ}]=\frac{1}{\sqrt{2\pi}}\int e^{-(x-t)^2/2+t^2/2}\,dx=e^{t^2/2}$, and $X=\mu+\sigma Z$.

Theorem 7.8

(a) If X,Y are independent, $M_{X+Y}(t)=M_X(t)M_Y(t)$. (b) (Uniqueness) If $M_X=M_Y$ on a neighbourhood of 0, then X and Y have the same distribution.

(a) is Proposition 7.2 with $g=e^{tx}$, $h=e^{ty}$; (b) is quoted without proof. Together they make Proposition 6.2 immediate: e.g. $e^{\mu_1 t+\sigma_1^2t^2/2}e^{\mu_2 t+\sigma_2^2t^2/2}$ is the normal mgf with parameters $\mu_1+\mu_2$, $\sigma_1^2+\sigma_2^2$.

Example 7.10

(a) χ^2_n: $E[e^{tZ^2}]=(1-2t)^{-1/2}$, so $Z_1^2+\cdots+Z_n^2$ has mgf $(1-2t)^{-n/2}$, the Gamma$(n/2,1/2)$ mgf.
(b) Random sums: $E[e^{tS}\mid N]=M_X(t)^N$, so $M_S(t)=E[M_X(t)^N]$. Differentiating at 0 recovers $E[S]=E[N]\mu$ and the variance formula of Example 7.9.
(c) If $M(t)=e^{3(e^t-1)}$, then $X\sim$ Poisson(3) and $P(X=0)=e^{-3}$.

The **joint mgf** $M(t_1,\dots,t_n)=E[e^{t_1X_1+\cdots+t_nX_n}]$ determines the joint law, and $X_1,\dots,X_n$ are independent iff it factors into the product of the marginal mgfs.

Example 7.11

Let X,Y be i.i.d. $N(\mu,\sigma^2)$. Then $X+Y$ and $X-Y$ are independent: $t(X+Y)+s(X-Y)=(t+s)X+(t-s)Y$, so

$$E[e^{t(X+Y)+s(X-Y)}]=e^{2\mu t+\sigma^2(t^2+s^2)}=E[e^{t(X+Y)}]\,E[e^{s(X-Y)}].$$

7.6 The multivariate normal distribution

If $Z_1,\dots,Z_n$ are independent $N(0,1)$ and $X_i=\mu_i+\sum_j a_{ij}Z_j$ $(i=1,\dots,m)$, then $(X_1,\dots,X_m)$ is **multivariate normal**. Every linear combination $\sum t_iX_i$ is then normal, with mean $\sum t_i\mu_i$ and variance $\sum_{i,j}t_it_j\operatorname{Cov}(X_i,X_j)$, so

$$M(t_1,\dots,t_m)=\exp\Big\{\sum_i t_i\mu_i+\frac{1}{2}\sum_{i,j}t_it_j\operatorname{Cov}(X_i,X_j)\Big\}:$$

the joint law is determined by the means and covariances. In particular, uncorrelated components of a multivariate normal vector are independent (the mgf factors).

Theorem 7.9 (Sample mean and variance)

If $X_1, \ldots, X_n$ are i.i.d. $N(\mu, \sigma^2)$, then $\bar{X}$ and S^2 are independent, $\bar{X} \sim N(\mu, \sigma^2/n)$, and $(n-1)S^2/\sigma^2 \sim \chi^2_{n-1}$.

Proof. $(\bar{X}, X_1 - \bar{X}, \ldots, X_n - \bar{X})$ is multivariate normal, and $\operatorname{Cov}(\bar{X}, X_i - \bar{X}) = \sigma^2/n - \sigma^2/n = 0$, so $\bar{X}$ is independent of the deviations and hence of S^2. Next,

$$\sum_{i=1}^{n} \Big(\frac{X_i - \mu}{\sigma}\Big)^2 = \frac{(n-1)S^2}{\sigma^2} + \Big(\frac{\bar{X} - \mu}{\sigma/\sqrt{n}}\Big)^2.$$

The left side is χ^2_n and the last term is χ^2_1 independent of the first; dividing mgfs, $(n-1)S^2/\sigma^2$ has mgf $(1-2t)^{-n/2}/(1-2t)^{-1/2} = (1-2t)^{-(n-1)/2}$. ■

7.7 General definition of expectation

For a nonnegative X that is neither discrete nor continuous one defines $E[X]$ as the limit of $E[X_n]$ for simple $X_n \uparrow X$ (the Lebesgue integral), or as $\int_0^\infty P(X > x)\,dx$; for general X, $E[X] = E[X^+] - E[X^-]$. All properties above continue to hold. If $X = Y$ with probability p and W otherwise, $E[X] = pE[Y] + (1-p)E[W]$.

Lecture 8. Limit Theorems

8.1 Markov and Chebyshev inequalities; the weak law

Proposition 8.1 (Markov's inequality)

If $X \geq 0$ and $a > 0$, then $P(X \geq a) \leq E[X]/a$.

Proof. $a\mathbb{1}\{X \geq a\} \leq X$; take expectations. ∎

Proposition 8.2 (Chebyshev's inequality)

If X has mean μ and variance σ^2, then $P(|X - \mu| \geq k) \leq \sigma^2/k^2$ for $k > 0$.

Proof. Apply Markov's inequality to $(X - \mu)^2$ and $a = k^2$. ∎

These bounds use only the mean (and variance), so they are often crude, but they hold for every distribution.

Example 8.1

(a) Weekly output has mean 50. Then $P(X \geq 75) \leq 50/75 = 2/3$. If also $\mathrm{Var}(X) = 25$, $P(40 < X < 60) \geq 1 - 25/100 = 3/4$.
(b) For $X \sim U(0, 10)$, Chebyshev gives $P(|X - 5| > 4) \leq (25/3)/16 \approx 0.52$, whereas the true value is 0.2.
(c) If $\mathrm{Var}(X) = 0$, then $P(|X - \mu| > 1/n) = 0$ for all n, and by continuity $P(X = \mu) = 1$.

Theorem 8.3 (Weak law of large numbers)

If $X_1, X_2, \ldots$ are i.i.d. with mean μ, then for every $\epsilon > 0$, $P(|\bar{X}_n - \mu| \geq \epsilon) \to 0$ as $n \to \infty$.

Proof. Assuming finite variance σ^2: $E[\bar{X}_n] = \mu$, $\mathrm{Var}(\bar{X}_n) = \sigma^2/n$, and Chebyshev gives the bound $\sigma^2/(n\epsilon^2) \to 0$. ∎

8.2 The central limit theorem

Theorem 8.4 (Central limit theorem)

If $X_1, X_2, \ldots$ are i.i.d. with mean μ and variance $\sigma^2 \in (0, \infty)$, then for every a,

$$P\Big(\frac{X_1 + \cdots + X_n - n\mu}{\sigma\sqrt{n}} \leq a\Big) \to \Phi(a) \qquad (n \to \infty).$$

The key tool is: if $M_{Z_n}(t) \to M_Z(t)$ for all t, then $F_{Z_n}(a) \to F_Z(a)$ at continuity points of F_Z (quoted).

Proof sketch. Take $\mu = 0$, $\sigma = 1$, and assume the mgf M of X_i exists near 0. Let $L(t) = \log M(t)$; then $L(0) = 0$, $L'(0) = \mu = 0$, $L''(0) = E[X^2] - \mu^2 = 1$. The mgf of $Z_n = \sum X_i/\sqrt{n}$ is $M(t/\sqrt{n})^n$, and by L'Hôpital twice (in $x = 1/\sqrt{n}$),

$$nL\Big(\frac{t}{\sqrt{n}}\Big) = \frac{L(tx)}{x^2} \to \frac{t^2 L''(0)}{2} = \frac{t^2}{2}.$$

So $M_{Z_n}(t) \to e^{t^2/2}$, the $N(0,1)$ mgf. (Without an mgf the same argument works with characteristic functions.) ∎

The theorem explains why measurement errors, heights, and sums of many small effects are approximately normal. It also holds for independent, non-identically distributed X_i that are uniformly bounded with $\sum \mathrm{Var}(X_i) = \infty$: then $(\sum(X_i - \mu_i))/\sqrt{\sum \sigma_i^2}$ is approximately $N(0,1)$.

Example 8.2 (How many measurements?)

An astronomer averages n independent measurements with variance 4. To be 95% sure that the average is within 0.5 of the true value, the CLT requires $P(|Z| \le 0.5\sqrt{n}/2) \approx 0.95$, i.e. $0.25\sqrt{n} \ge 1.96$, $n \ge 62$. Chebyshev's inequality would demand $4/(0.25n) \le 0.05$, i.e. $n \ge 320$: guaranteed for every distribution, but far more conservative.

Example 8.3

(a) The sum S of 10 fair dice has mean 35 and variance $350/12$; with the continuity correction, $P(30 \le S \le 40) \approx 2\Phi(5.5/5.40) - 1 \approx 0.69$.
(b) The sum of 10 independent $U(0,1)$ variables has mean 5 and variance $10/12$, so $P(\sum X_i > 6) \approx 1 - \Phi(1.10) \approx 0.137$.
(c) Exam grading times are i.i.d. with mean 20 and standard deviation 4 minutes. The probability that the first 25 exams are graded within 450 minutes is approximately $\Phi\big((450 - 500)/20\big) = \Phi(-2.5) \approx 0.006$.

8.3 The strong law of large numbers

Theorem 8.5 (Strong law of large numbers)

If $X_1, X_2, \ldots$ are i.i.d. with finite mean μ, then $P\big(\lim_{n\to\infty} \bar{X}_n = \mu\big) = 1$.

Proof. Under the extra assumption $K = E[X_i^4] < \infty$, with $\mu = 0$. Expanding $E[S_n^4]$, the terms with some index appearing exactly once vanish, leaving $nE[X^4] + 6\binom{n}{2}E[X_i^2 X_j^2] \le nK + 3n(n-1)K$ (as $E[X^2]^2 \le E[X^4]$). Hence $E\big[\sum_n S_n^4/n^4\big] \le \sum_n (K/n^3 + 3K/n^2) < \infty$, so $\sum_n S_n^4/n^4 < \infty$ with probability 1, which forces $S_n^4/n^4 \to 0$, i.e. $\bar{X}_n \to 0$. ∎

The weak law says that for each large n, $\bar{X}_n$ is *probably* close to μ; the strong law says that, with probability 1, the whole sequence eventually stays close to μ. Applied to $X_i = \mathbb{1}\{E \text{ occurs on trial } i\}$, it justifies the frequency interpretation: the long-run proportion of trials on which E occurs equals $P(E)$ with probability 1.

8.4 Other inequalities

Proposition 8.6 (One-sided Chebyshev)

If $E[X] = 0$ and $\mathrm{Var}(X) = \sigma^2$, then $P(X \ge a) \le \sigma^2/(\sigma^2 + a^2)$ for $a > 0$.

Proof. For $b > 0$, $P(X \ge a) \le P\big((X+b)^2 \ge (a+b)^2\big) \le (\sigma^2 + b^2)/(a+b)^2$ by Markov; choose $b = \sigma^2/a$. ∎

Proposition 8.7 (Chernoff bounds)

If $M(t) = E[e^{tX}]$, then $P(X \geq a) \leq e^{-ta}M(t)$ for all $t > 0$ and $P(X \leq a) \leq e^{-ta}M(t)$ for all $t < 0$.

Proof. For $t > 0$, $P(X \geq a) = P(e^{tX} \geq e^{ta}) \leq E[e^{tX}]e^{-ta}$ by Markov. ■

One then minimizes over t.

Example 8.4

(a) Standard normal: $e^{-ta+t^2/2}$ is minimized at $t = a$, so $P(Z \geq a) \leq e^{-a^2/2}$ for $a > 0$.
(b) Poisson(λ) and $i > \lambda$: minimizing $e^{\lambda(e^t-1)-ti}$ at $e^t = i/\lambda$ gives $P(X \geq i) \leq e^{-\lambda}(e\lambda)^i/i^i$.
(c) A gambler wins or loses 1 with probability $\frac{1}{2}$ each play. After n plays, S_n has $M(t) = \cosh(t)^n \leq e^{nt^2/2}$, so $P(S_n \geq a) \leq e^{-a^2/(2n)}$ (take $t = a/n$); e.g. $P(S_{10} \geq 6) \leq e^{-1.8} \approx 0.165$, while the exact value is 0.0547.

Proposition 8.8 (Jensen's inequality)

If f is convex, then $E[f(X)] \geq f(E[X])$ whenever the expectations exist.

Proof. With $\mu = E[X]$, convexity gives $f(x) \geq f(\mu) + f'(\mu)(x - \mu)$ (for twice differentiable f, by Taylor's theorem with $f'' \geq 0$). Take expectations. ■

For example, an investor choosing between a sure return μ and a random return X with mean μ, with concave utility u, prefers the sure amount since $E[u(X)] \leq u(\mu)$.

8.5 The Poisson approximation, quantified

Let $X = \sum_{i=1}^n X_i$ with independent $X_i \sim$ Bernoulli(p_i), and $Y \sim$ Poisson(λ) with $\lambda = \sum p_i$. Coupling each X_i with a Poisson(p_i) variable Y_i so that they agree with probability at least $1 - p_i^2$, one obtains, for every set A of integers,

$$|P(X \in A) - P(Y \in A)| \leq \sum_{i=1}^{n} p_i^2.$$

For Bin(n, p) the bound is $np^2 = \lambda p$: the approximation is good when p is small, whatever n is.

Example 8.5 (Birthdays, again)

Let $X_1, X_2, \ldots$ be the birthdays of people arriving in sequence and N the number needed for the first shared birthday. For n people, the event $N > n$ means that none of the trials "person j shares with one of persons $1, \ldots, j-1$" succeeds; these have success probabilities $(j-1)/365$ given no earlier success, and an exact product gives $P(N > n) = \prod_{j=1}^n (1 - \frac{j-1}{365}) \approx \exp\{-\sum_j \frac{j-1}{365}\} = e^{-n(n-1)/730}$, the Poisson-paradigm answer of Example 4.6.