Lecture 3. Cholesky Decomposition

Existence and uniqueness, algorithm and cost, pivoting, applications, sweep operator

Contents

3.1 The Cholesky decomposition **2**

3.2 Applications **3**

3.3 The sweep operator **4**

Mantra 2

The structure should be exploited whenever possible in solving a problem.

Common structures are symmetry, positive (semi)definiteness, sparsity, Kronecker products and low rank. Statisticians mostly meet positive (semi)definite matrices: covariance matrices and Gram matrices $X^\top X$. A matrix M is positive definite (semidefinite) if $x^\top Mx > 0$ (≥ 0) for all $x \neq 0$; since $x^\top Mx = x^\top \frac{M+M^\top}{2} x$, we may assume M symmetric.

3.1 The Cholesky decomposition

Theorem 3.1 (Cholesky)

If $A \in \mathbb{R}^{n\times n}$ is symmetric positive definite, then $A = LL^\top$ for a unique lower triangular L with positive diagonal entries.

Proof. Induction on n. For $n = 1$, $a = \sqrt{a}\sqrt{a}$ with $a > 0$. For $n > 1$, partition and match blocks:

$$\begin{pmatrix} a_{11} & \mathbf{a}^\top \\ \mathbf{a} & A_{22} \end{pmatrix} = \begin{pmatrix} \ell_{11} & \mathbf{0}^\top \\ \mathbf{b} & L_{22} \end{pmatrix} \begin{pmatrix} \ell_{11} & \mathbf{b}^\top \\ \mathbf{0} & L_{22}^\top \end{pmatrix} \iff a_{11} = \ell_{11}^2, \quad \mathbf{a} = \ell_{11}\mathbf{b}, \quad A_{22} = \mathbf{b}\mathbf{b}^\top + L_{22}L_{22}^\top.$$

Since $a_{11} = e_1^\top Ae_1 > 0$, the conditions force $\ell_{11} = \sqrt{a_{11}}$ and $\mathbf{b} = \mathbf{a}/\sqrt{a_{11}}$. (Every principal submatrix of a positive definite matrix is positive definite, e.g. $\mathbf{y}^\top A_{22}\mathbf{y} = (0, \mathbf{y}^\top)A(0, \mathbf{y}^\top)^\top > 0$; but the recursion needs more.) The **Schur complement** $S = A_{22} - \mathbf{b}\mathbf{b}^\top = A_{22} - \mathbf{a}\mathbf{a}^\top/a_{11}$ is positive definite: for $\mathbf{y} \neq \mathbf{0}$, put $x = (-\mathbf{a}^\top\mathbf{y}/a_{11}, \mathbf{y}^\top)^\top \neq 0$; then

$$0 < x^\top Ax = \frac{(\mathbf{a}^\top\mathbf{y})^2}{a_{11}} - 2\frac{(\mathbf{a}^\top\mathbf{y})^2}{a_{11}} + \mathbf{y}^\top A_{22}\mathbf{y} = \mathbf{y}^\top S\mathbf{y}.$$

By the induction hypothesis $S = L_{22}L_{22}^\top$ for a unique such L_{22}, which completes existence and uniqueness. ■

The proof is the algorithm: compute the first column of L, update the Schur complement, and recurse; this can be done in place in the lower triangle of A.

Cost. The Schur complement updates are symmetric, so only half the entries are computed:

$$\tfrac{1}{2}\{2(n-1)^2 + 2(n-2)^2 + \cdots + 2\} \approx \tfrac{1}{3}n^3 \text{ flops}$$

plus n square roots, *half* the cost of LU. The algorithm is very stable without pivoting, and it fails (a nonpositive number under a square root) exactly when A is not positive definite. So Cholesky is the cheapest way to *test* positive definiteness.

Example 3.1

For

$$A = \begin{pmatrix} 4 & 6 & 4 & 4 \\ 6 & 37/4 & 11/2 & 6 \\ 4 & 11/2 & 29/4 & 7 \\ 4 & 6 & 7 & 33/4 \end{pmatrix}$$

the first column of L is $(4, 6, 4, 4)^\top/2 = (2, 3, 2, 2)^\top$. The Schur complement $A_{22} - \mathbf{b}\mathbf{b}^\top$ has first column $(37/4 - 9,\ 11/2 - 6,\ 6 - 6)^\top = (1/4, -1/2, 0)^\top$, so $\ell_{22} = 1/2$ and the second

column of L is $(1/2, -1, 0)^\top$. Continuing, the pivots are $4, \frac{1}{4}, \frac{9}{4}, \frac{1}{4}$, all positive, and

$$L = \begin{pmatrix} 2 & 0 & 0 & 0 \\ 3 & 1/2 & 0 & 0 \\ 2 & -1 & 3/2 & 0 \\ 2 & 0 & 2 & 1/2 \end{pmatrix}.$$

Since the algorithm ran to completion with positive pivots, $A = LL^\top$ is positive definite (and $\det A = (2 \cdot \frac{1}{2} \cdot \frac{3}{2} \cdot \frac{1}{2})^2 = \frac{9}{16}$).

Pivoting for semidefinite matrices. If A is only positive semidefinite (e.g. $X^\top X$ with X rank deficient), a zero pivot can occur. With **symmetric pivoting**, at stage k the largest remaining diagonal entry is moved (by a symmetric row and column swap) to the pivot position. If the largest remaining diagonal entry is 0, the whole remaining block is 0: a positive semidefinite matrix with zero diagonal entry a_{ii} has zero row i, since $|a_{ij}| \le \sqrt{a_{ii}a_{jj}}$. The algorithm then stops, giving

$$PAP^\top = LL^\top, \qquad L \in \mathbb{R}^{n \times r}, \quad r = \operatorname{rank} A.$$

LAPACK provides `potrf` (no pivoting) and `pstrf` (pivoting); R's `chol` returns the upper factor $R = L^\top$, $A = R^\top R$.

3.2 Applications

Multivariate normal log-likelihood. For $y \sim N_n(0, \Sigma)$,

$$\log f(y) = -\frac{n}{2}\log(2\pi) - \frac{1}{2}\log\det\Sigma - \frac{1}{2}y^\top\Sigma^{-1}y.$$

- Method 1: invert Σ ($2n^3$), form the quadratic form ($2n^2$), compute the determinant by LU ($\frac{2}{3}n^3$).
- Method 2: factor $\Sigma = LL^\top$ ($\frac{1}{3}n^3$), solve $Lx = y$ by forward substitution (n^2), then $y^\top\Sigma^{-1}y = x^\top x$ ($2n$) and $\log\det\Sigma = 2\sum_i \log \ell_{ii}$ (n).

Method 2 is about 8 times cheaper and more accurate, and the log-determinant computed this way cannot overflow. To evaluate the density at many points, factor once and pay n^2 per point.

Linear regression. For $X \in \mathbb{R}^{n \times p}$, solve the normal equations $X^\top X\beta = X^\top y$ by Cholesky:

form $X^\top X$	np^2 (symmetric)
form $X^\top y$	$2np$
Cholesky of $X^\top X$	$\frac{1}{3}p^3$
two triangular solves	$2p^2$
$(X^\top X)^{-1}$ for standard errors	$\frac{4}{3}p^3$

In total about $np^2 + \frac{1}{3}p^3$ flops, or $np^2 + \frac{5}{3}p^3$ with standard errors. This is the fastest method when $n \gg p$, but forming $X^\top X$ squares the condition number (Lecture 4).

3.3 The sweep operator

For a symmetric $A = (a_{ij}) \in \mathbb{R}^{n\times n}$ with $a_{kk} \neq 0$, **sweeping** on the kth diagonal entry maps A to $\hat{A} = (\hat{a}_{ij})$ with

$$\hat{a}_{kk} = -\frac{1}{a_{kk}}, \qquad \hat{a}_{ik} = \frac{a_{ik}}{a_{kk}}, \qquad \hat{a}_{kj} = \frac{a_{kj}}{a_{kk}}, \qquad \hat{a}_{ij} = a_{ij} - \frac{a_{ik}a_{kj}}{a_{kk}} \quad (i, j \neq k).$$

The last formula is one step of Gaussian elimination (a Schur complement update), so a sweep costs about n^2 flops ($n^2/2$ using symmetry), and it can be done in place. The **inverse sweep** on entry k is

$$\check{a}_{kk} = -\frac{1}{a_{kk}}, \qquad \check{a}_{ik} = -\frac{a_{ik}}{a_{kk}}, \qquad \check{a}_{kj} = -\frac{a_{kj}}{a_{kk}}, \qquad \check{a}_{ij} = a_{ij} - \frac{a_{ik}a_{kj}}{a_{kk}};$$

substituting $\hat{A}$ shows directly that it restores A (e.g. $\hat{a}_{ij} - \hat{a}_{ik}\hat{a}_{kj}/\hat{a}_{kk} = a_{ij} - a_{ik}a_{kj}/a_{kk} + a_{ik}a_{kj}/a_{kk} = a_{ij}$).

Lemma 3.2 (Exchange property)

Let $U = [u_1, \ldots, u_n]$ and $V = [v_1, \ldots, v_n]$ be $p \times n$ with $V = UA$. After sweeping on entry k,

$$[v_1, \ldots, v_{k-1}, -u_k, v_{k+1}, \ldots, v_n] = [u_1, \ldots, u_{k-1}, v_k, u_{k+1}, \ldots, u_n]\,\hat{A}.$$

That is, sweeping exchanges the roles of u_k and v_k (up to sign).

Proof. Column k of the right side is $-v_k/a_{kk} + \sum_{i\neq k} u_i a_{ik}/a_{kk} = -(v_k - \sum_{i\neq k} u_i a_{ik})/a_{kk} = -u_k a_{kk}/a_{kk} = -u_k$, using $v_k = \sum_i u_i a_{ik}$. For $j \neq k$, column j is

$$v_k \frac{a_{kj}}{a_{kk}} + \sum_{i\neq k} u_i \Big(a_{ij} - \frac{a_{ik}a_{kj}}{a_{kk}}\Big) = \sum_{i\neq k} u_i a_{ij} + \frac{a_{kj}}{a_{kk}}\Big(v_k - \sum_{i\neq k} u_i a_{ik}\Big) = \sum_{i\neq k} u_i a_{ij} + a_{kj}u_k = v_j.$$

■

Proposition 3.3

Partition $A = \left(\begin{smallmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{smallmatrix}\right)$ with $A_{11} \in \mathbb{R}^{k\times k}$ invertible. Sweeping on the diagonal entries $1, \ldots, k$ (in any order, as long as each pivot is nonzero) gives

$$\hat{A} = \begin{pmatrix} -A_{11}^{-1} & A_{11}^{-1}A_{12} \\ A_{21}A_{11}^{-1} & A_{22} - A_{21}A_{11}^{-1}A_{12} \end{pmatrix}.$$

Proof. Apply the lemma to $A = IA$ ($U = I$, $V = A$) for each sweep. After sweeping $1, \ldots, k$,

$$\begin{pmatrix} -I_k & A_{12} \\ 0 & A_{22} \end{pmatrix} = \begin{pmatrix} A_{11} & 0 \\ A_{21} & I \end{pmatrix} \begin{pmatrix} \hat{A}_{11} & \hat{A}_{12} \\ \hat{A}_{21} & \hat{A}_{22} \end{pmatrix}.$$

Reading off blocks: $A_{11}\hat{A}_{11} = -I$, $A_{11}\hat{A}_{12} = A_{12}$, $A_{21}\hat{A}_{11} + \hat{A}_{21} = 0$ and $A_{21}\hat{A}_{12} + \hat{A}_{22} = A_{22}$, which give the four blocks. The result depends only on the set of swept indices, so sweeps commute. ■

Regression by sweeping. Sweep the $(p+1)\times(p+1)$ matrix $\left(\begin{smallmatrix} X^\top X & X^\top y \\ y^\top X & y^\top y \end{smallmatrix}\right)$ on its first p diagonal entries. By the proposition the result is

$$\begin{pmatrix} -(X^\top X)^{-1} & (X^\top X)^{-1}X^\top y \\ y^\top X(X^\top X)^{-1} & y^\top y - y^\top X(X^\top X)^{-1}X^\top y \end{pmatrix} = \begin{pmatrix} -\sigma^{-2}\operatorname{Cov}(\hat{\beta}) & \hat{\beta} \\ \hat{\beta}^\top & \|y - \hat{y}\|^2 \end{pmatrix}:$$

the coefficients, their (unscaled) covariance and the residual sum of squares appear at once. This makes sweeping ideal for **stepwise regression**. Adding variable j to the current model is one sweep on entry j and removing it is one inverse sweep, each $O(p^2)$, instead of refitting from scratch; the swept matrix always holds the current $\hat{\beta}$, standard errors and RSS, from which the F (or $t^2 = \hat{\beta}_j^2 / \widehat{\mathrm{Var}}\hat{\beta}_j$) statistics for entering or leaving variables are read off. A pivot that is (numerically) zero signals a variable that is collinear with those already in the model. The price is that it works with $X^\top X$, with the conditioning of the normal equations.